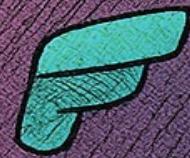


FebCon Vegas 2075



Welcome
to FabCon

#FABCON



FebCon Vegas 2015



Welcome
to FabCon

#FABCON

News and Recap

Andrea Benedetti

OneSecurity

Intro

- OneLake Security evolves the security model that already exists in Microsoft Fabric, adding the ability to define granular security rules (at the row and column level) directly within OneLake.
- This makes security simpler, more consistent, and easier to manage, so you don't have to redefine specific rules in each individual analytics engine like Spark, SQL Endpoint, and Power BI.

Main changes

- Create and manage centralized security roles directly in OneLake.
- Ability to granularly define:
 - Data access, by selecting specific tables or folders.
 - Row and column-level security to precisely limit access to sensitive data.
- New user interface to manage roles, clearly view assigned data, and quickly change user memberships.

All Lakehouse roles > NewRole

NewRole

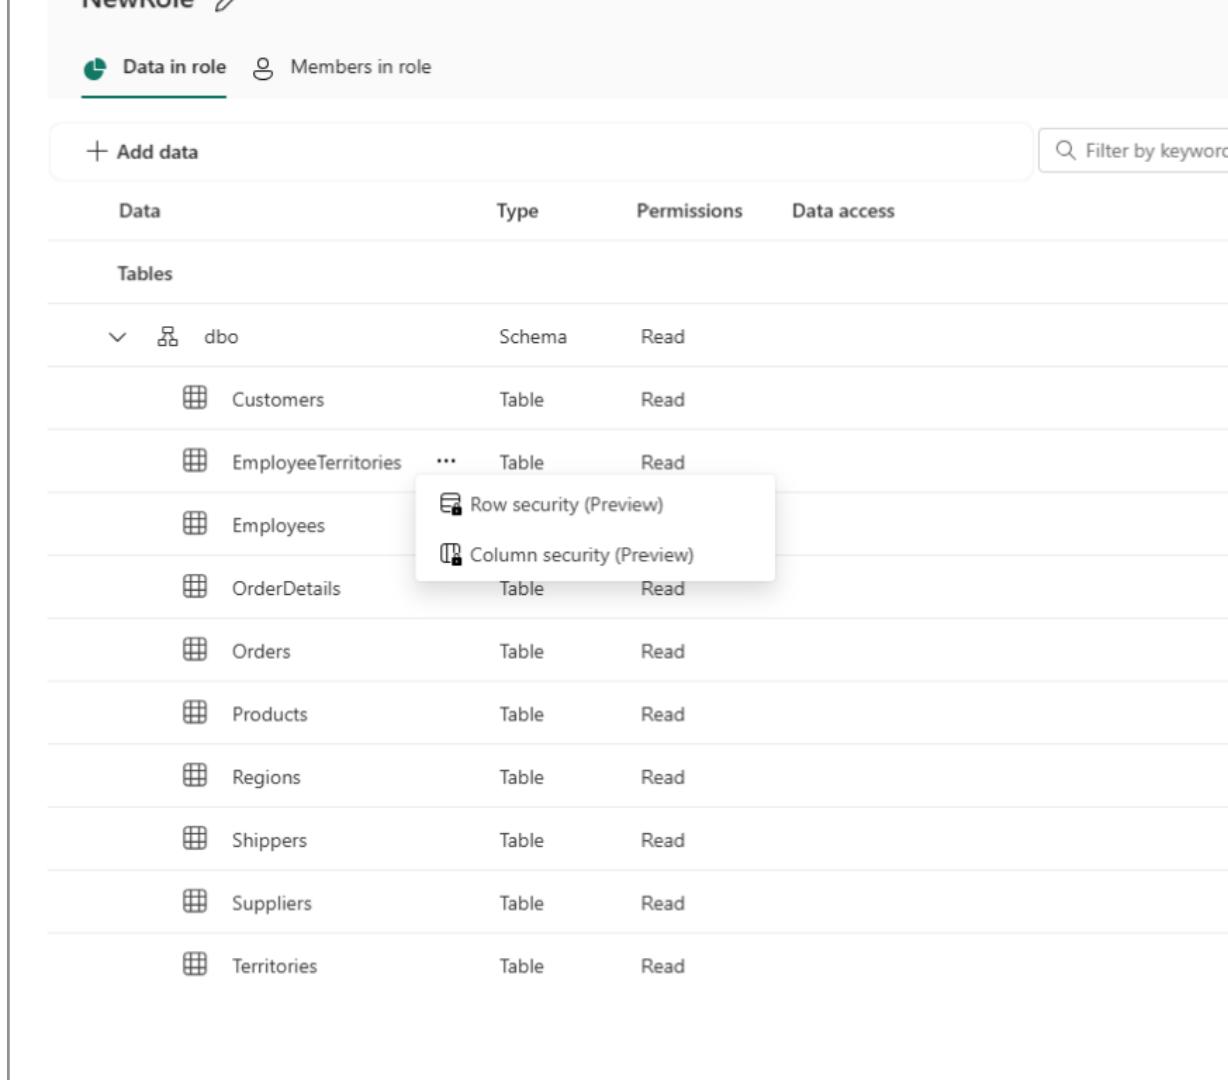
Data in role Members in role

+ Add data Filter by keyword

| Data | Type | Permissions | Data access |
|---------------------|--------|-------------|-------------|
| Tables | | | |
| dbo | Schema | Read | |
| Customers | Table | Read | |
| EmployeeTerritories | Table | Read | |
| Employees | Table | Read | |
| OrderDetails | Table | Read | |
| Orders | Table | Read | |
| Products | Table | Read | |
| Regions | Table | Read | |
| Shippers | Table | Read | |
| Suppliers | Table | Read | |
| Territories | Table | Read | |

Row security (Preview)

Column security (Preview)



Integrated support with Fabric tools

Spark

- Automatically supports rules defined in OneLake Security in Spark notebooks, including row and column-level restrictions.

SQL Analytics Endpoint

- Introduction of the "user identity" mode that allows SQL endpoints to directly comply with OneLake security rules.
- New endpoints will automatically adopt this mode, while existing ones can be easily converted.
- In "user identity" mode, all security rules are managed exclusively in OneLake.

Semantic Models (Power BI in modalità Direct Lake)

- Automatically enforce OneLake Security rules, ensuring that users only see the data they actually have access to when creating reports or editing models.

Link

- OneLake security – [early access form](#)
- Documentation
 - [OneLake Security Overview](#)
- Video
 - [OneLake Security Demo](#)

Empowering agentic AI by
integrating Fabric data agents
with Azure AI Foundry

Agentic AI integrating Fabric with AI Foundry

- Greater integration between Microsoft Fabric and Azure AI Foundry, enabling the creation of custom conversational AI agents that can operate autonomously and generate accurate, relevant, and data-driven responses.

Microsoft Fabric Data Agents: What are they?

- Intelligent assistants capable of instantly analyzing structured and semantic data.
- They provide a simple and effective conversational experience to access business data from multiple sources.
- They enable users to quickly create data-driven conversational experiences by adding advanced analytical capabilities to their AI agents.

Main news

- AI agents integrated with business data
 - AI agents, called "Fabric data agents", can use different data sources such as lakehouses, warehouses, Power BI semantic models and KQL databases, accessing information with specialized query languages (SQL, KQL, DAX).
 - These agents are able to decide for themselves what data to use, how to combine it, and which insights are most relevant.
- Connecting Fabric to Azure AI Foundry
 - The integration combines the advanced analytics and data management capabilities of Fabric with the generative (GenAI) technology of Azure AI Foundry.
 - It enables the development of conversational AI agents that integrate both unstructured data (e.g., SharePoint, Azure AI Search) and structured and semantic data from Microsoft OneLake.

OneLake's central role

- OneLake is the unified data lake that centralizes all business data, regardless of where it comes from (internal or external, such as Snowflake or other databases).
- It provides a solid and comprehensive foundation of knowledge for AI agents, ensuring analytics based on a comprehensive view of the organization's data.

Built-in security

- Identity Passthrough / On-Behalf-Of - OBO
- The integration supports end-user identity (OBO)-based authentication, ensuring that each user receives AI responses based solely on the data they have access to.
- It provides native and secure control over the management and access to sensitive data.

Fabric Demo

The screenshot shows the Azure AI Foundry interface for the AdventureWorksAIProject. The top navigation bar includes 'Azure AI Foundry' (with a blue icon), 'AdventureWorksAIProject' (with a dropdown arrow), 'Agents' (selected), 'All hubs + projects' (with a dropdown arrow), and a user profile for 'Project contosoaiproject'. A green success message at the top states: 'Success: Agent with id asst_XgeKSy42n2lbw6Duhi81JTZw created successfully'.

The left sidebar has a 'Build and customize' section expanded, showing 'Agents PREVIEW' selected. Other sections include 'Templates PREVIEW', 'Fine-tuning', 'Prompt flow', 'Assess and improve' (with 'Tracing PREVIEW'), 'Evaluation', and 'Safety + security'. Under 'My assets', there are 'Models + endpoints', 'Data + indexes', and 'Web apps'. At the bottom is a 'Management center' button.

The main content area is titled 'Create and debug your agents'. It features a table with columns 'Name', 'ID', and 'Created'. One row is listed: 'Agent551' (ID: asst_XgeKSy42n2lbw6Duhi81JTZw, Created: Mar 27, 2024). Below the table are buttons for '+ New agent', 'Edit connected resources', 'Refresh', 'Delete', and 'Copy'. A search bar and a 'Filter' icon are also present.

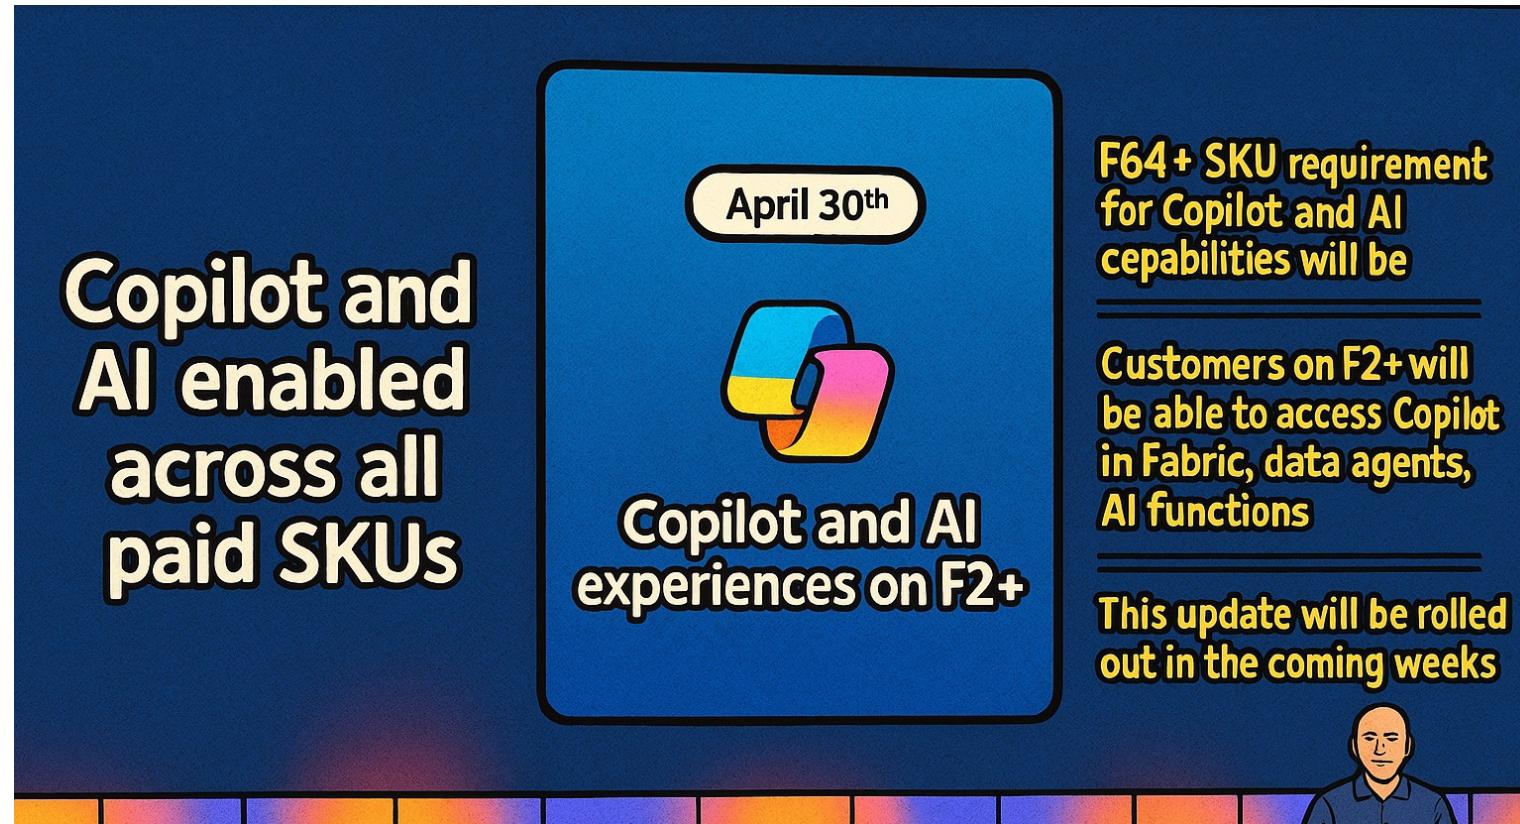
To the right, a large panel titled 'Setup' contains fields for 'Agent id' (set to 'asst_XgeKSy42n2lbw6Duhi81JTZw'), 'Agent name' (set to 'Agent551'), 'Azure AI Services or Azure OpenAI resource' (set to 'zhuoqunliaiservice'), 'Deployment' (set to 'gpt-4o (version:2024-08-06)'), and an 'Instructions' text area. A 'Try in playground' button is located above the deployment section.

- Video
 - [Integrating Fabric with Azure AI Foundry Demo](#)
- Link
 - [Expand Azure AI Agent with New Knowledge Tools: Microsoft Fabric and Tripadvisor | Microsoft Community Hub](#)
 - [How to use the data agents in Microsoft Fabric with Azure AI Agent Service - Azure OpenAI | Microsoft Learn](#)

Copilot and AI capabilities in Fabric

Universal access to Copilot and AI capabilities in Microsoft Fabric

- Elimination of SKU constraint for Copilot / AI usage
- In the coming weeks, all F2+ licensed users will be able to use Copilot, Fabric Data Agents, and other advanced AI features, at no additional cost or configuration.



What does this mean concretely?

- AI available to everyone
 - Whether you're an analyst, data scientist, or manager, you can now easily leverage advanced AI tools built right into the Fabric platform.
- Increased productivity
 - Copilot makes everyday tasks like creating data visualizations in Power BI or automating complex analytic tasks quick and easy.
- Fabric Data Agents
 - These agents allow for simple and natural interaction with business data. They understand the specific context of your organization's data, dramatically reducing the time it takes to generate relevant insights.

Why is this important?

- Democratization of AI
 - Microsoft intends to make advanced AI tools accessible to an increasing number of users and organizations.
- Ease and flexibility of use
 - The introduction of the new Fabric Copilot capability further simplifies management, reduces costs, and increases overall efficiency in the adoption of AI technologies.
- Growth and innovation for all
 - Removing barriers to entry allows companies of all sizes to harness the transformative potential of AI, focusing on their business challenges and creating new opportunities for growth.

Link

- [Overview of Copilot in Fabric - Microsoft Fabric | Microsoft Learn](#)
- [Enable Copilot in Fabric - Microsoft Fabric | Microsoft Learn](#)

**Seamlessly migrate your data to
Fabric**

Migration Assistant da Azure Synapse Analytics verso Microsoft Fabric

- Microsoft has announced a new native tool in Fabric, called the Migration Assistant for Fabric Data Warehouse, to facilitate migration from Azure Synapse Analytics (Data Warehouse) to Microsoft Fabric.
- This feature will be available globally by the second week of April 2025.

Key benefits of the new migration experience

- Simplified migration
 - The assistant allows you to easily migrate metadata and data from your existing data warehouse in Azure Synapse Analytics to your Fabric Data Warehouse.
- Automatic schema conversion
 - The original schema is automatically adapted to the format supported by Fabric, including automatic conversion of data types and T-SQL scripts.
- AI support via Copilot
 - Copilot assists users in the automated resolution of migration errors, indicating how to correct any incompatibilities in the T-SQL script.

Four-step migration process

- Metadata migration
 - You use a DACPAC file exported from the source system to import objects such as tables, views, functions, stored procedures, and security objects.
 - The assistant automatically translates the object definitions so that they are compatible with Fabric.
- Troubleshooting
 - Any errors in the automatic migration are clearly indicated by the assistant.
 - Users can manually fix these errors, or use Copilot for AI-assisted fixes, which also explains the changes made.
- Copying data
 - A built-in wizard helps you quickly select tables and columns to migrate, define data mapping, and choose how to copy (full or incremental copies).
- Reconfiguring connections
 - Finally, ETL and reporting tools are reconnected to the new Data Warehouse Fabric.

More improvements coming soon

- Live connectivity with Azure Synapse and other data sources
- Enhanced Copilot with improved experiences, original and modified script comparison, and new auto-correction options
- Ability to view or restore the original migrated code
- Full-screen experience for the migration assistant
- Even smoother and more complete data copy process
- Automated and assisted management of connection reconfiguration
- Continuous investment to improve usability and SQL compatibility

Demo

The screenshot shows a Microsoft DXT Power BI workspace titled "Migration demo". The left sidebar includes icons for Home, Copilot, Create, Browse, OneLake, Apps, Metrics, Monitor, Learn, Real-Time, Workloads, Workspaces, Migration demo (which is selected), and copyjob1. The main area displays a large circular icon with two overlapping squares. Below it, the text reads "Choose from predesigned task flows or add a task to build one (preview)". A sub-instruction says "Select from one of Microsoft's predesigned task flows or add a task to start building one yourself." There are two buttons: "Select a predesigned task flow" and "Add a task". A table below lists three tasks:

| Name | Type | Task | Owner | Refreshed | Next refresh | Endorsement | Sensitivity | Included in app |
|----------------|-----------------------|------|----------------|-----------------------|--------------|-------------|-----------------------|-----------------|
| AdevntureWorks | Semantic model (d...) | — | Migration demo | 3/24/2025, 3:14:28 PM | N/A | — | Confidential\Micro... | ○ |
| AdventureWorks | Semantic model (d...) | — | Migration demo | 3/24/2025, 2:50:55 PM | N/A | — | Confidential\Micro... | ○ |
| copyjob1 | Copy job | — | Ancy Philip | — | — | — | Confidential\Micro... | ○ |

Mirroring

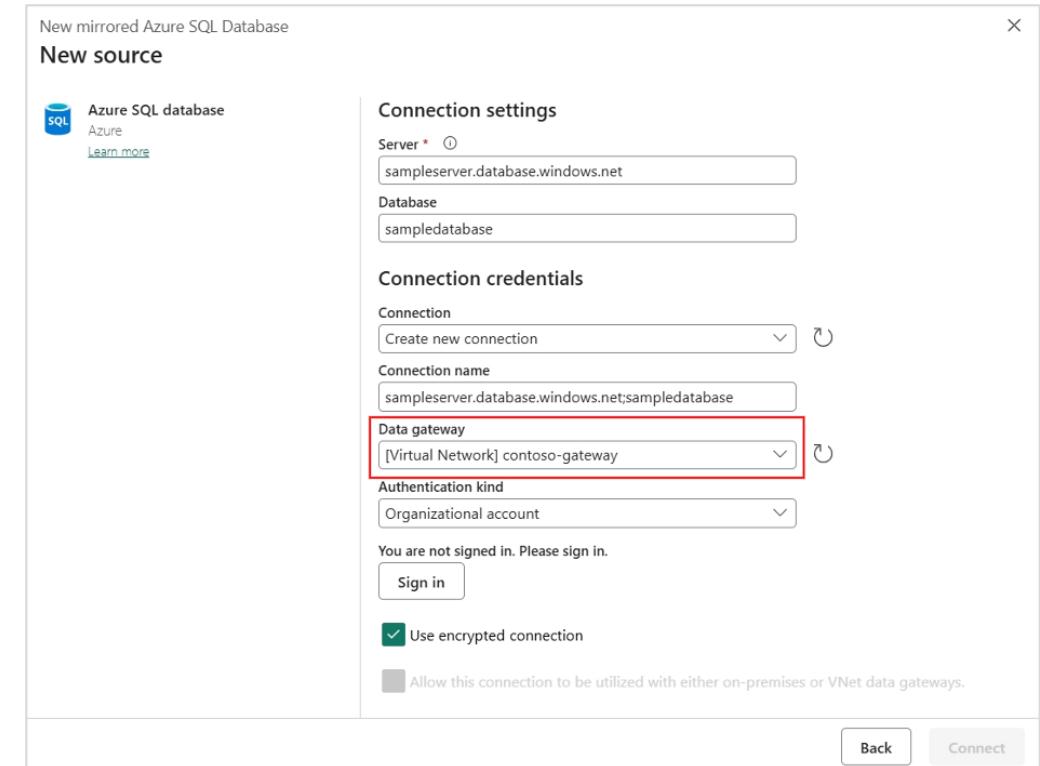
Intro

- Microsoft Fabric has announced important new features for the Mirroring service, which allows you to automatically replicate data from various sources to the unified OneLake data lake in real time.
- Mirroring allows you to keep your business data within Fabric up to date continuously, making it easily available for advanced analytics, machine learning, and reporting, without interfering with the original operational databases.

Key new features

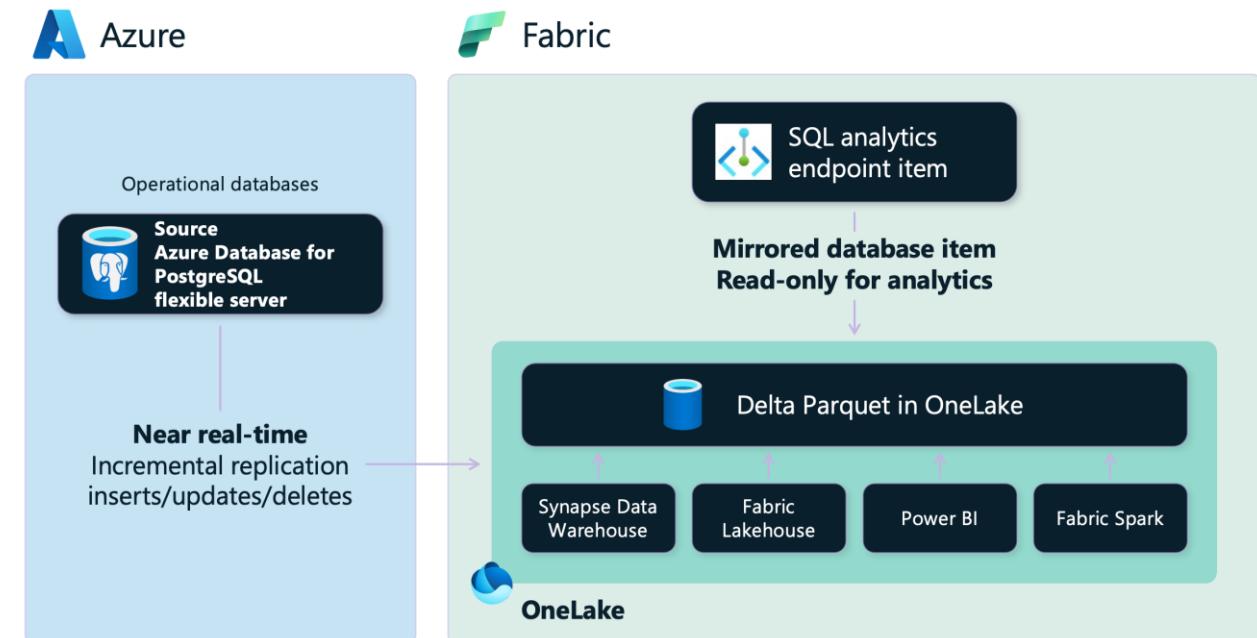
Database Support Behind Firewall

- Mirroring can now replicate data from databases behind firewalls or located on-premises, thanks to integration with:
 - **Azure SQL Database**
 - **Snowflake**
 - (soon also) **Azure SQL Managed Instance**
- Use **On-Premises Data Gateway** and **Virtual Network Data Gateway** for secure connections.



Mirroring per Azure Database for PostgreSQL Flexible Server

- Continuous, real-time replication from the PostgreSQL Flexible Server service to OneLake is now available.
- Allows you to use PostgreSQL data directly for reporting, analytics, and machine learning in Fabric.



Supporto CI/CD per Mirroring (GA)

- Mirroring ora supporta integrazioni CI/CD (Continuous Integration/Continuous Deployment).
- It can be integrated with Git and ALM Deployment pipelines to automate updates and configurations of replicated databases.

The screenshot shows the 'Mirroring Dev Workspace' interface. At the top, there are navigation links for 'Dev', 'View deployment pipeline', 'Create app', and 'Manage access'. Below the header is a toolbar with buttons for '+ New item', 'New folder', 'Import', 'Source control' (with a count of 0), 'Filter by keyword', 'Filter', and a search icon. The main area is a table listing database objects:

| Name | Git status | Type | Task | Owner |
|---------------------|------------|--------------------------|------|-------------------------|
| DemoAzureSQLDB-AWLT | Synced | Mirrored database | — | Linda Wang |
| DemoAzureSQLDB AWLT | — | Semantic model (default) | — | Mirroring Dev Workspace |
| DemoAzureSQLDB AWLT | — | SQL analytics endpoint | — | Linda Wang |
| SnowflakeSample | Synced | Mirrored database | — | Linda Wang |
| SnowflakeSample | — | Semantic model (default) | — | Mirroring Dev Workspace |
| SnowflakeSample | — | SQL analytics endpoint | — | Linda Wang |

Built-in monitoring in workspaces

- Detailed monitoring of the mirroring process is now available via **Workspace Monitoring**.
- It provides detailed operational logs, performance metrics, replication status, latencies, and possible errors.
- It allows direct KQL queries for point analysis, custom dashboards and setting specific alerts.

| Timestamp | OperationName | ItemId | ItemKind | ItemName | WorkspaceId |
|----------------------------|-------------------|--------------------------------------|------------------|------------------|---------------|
| > 2025-02-18 05:19:20.2140 | AddTable | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |
| > 2025-02-18 05:19:20.2160 | AddTable | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |
| > 2025-02-18 05:19:20.2180 | AddTable | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |
| > 2025-02-18 05:19:20.2190 | AddTable | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |
| > 2025-02-18 05:19:20.2210 | AddTable | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |
| > 2025-02-18 05:19:20.7880 | ReplicatingSchema | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |
| > 2025-02-18 05:19:20.7980 | ReplicatingSchema | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |
| > 2025-02-18 05:19:20.8140 | ReplicatingSchema | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |
| > 2025-02-18 05:19:20.8190 | ReplicatingSchema | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |
| > 2025-02-18 05:19:20.8610 | ReplicatingSchema | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |
| > 2025-02-18 05:19:35.8440 | ReplicatingSchema | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |
| > 2025-02-18 05:19:35.8520 | ReplicatingSchema | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |
| > 2025-02-18 05:19:35.8590 | ReplicatingSchema | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |
| > 2025-02-18 05:19:35.8790 | ReplicatingSchema | fca8cfcd-ff10-4e90-98b5-cdce/cacd0da | MirroredDatabase | SampleAzureSQLDB | 3023603b-43a1 |

Replicating Database Schemas

- Mirroring now also faithfully replicates the original schema structure of the source databases.
- It ensures consistency between original data and data replicated in OneLake, making it easy to use with Spark, SQL, and semantic models.

Delta Column Mapping Support

- You can now replicate columns with names that contain spaces or special characters with support for the Delta format.
- It makes it easier and more flexible to replicate complex or specially named data.

Open Mirroring: UX Enhancements and CSV Support

- The Open Mirroring feature, based on the open Delta Lake format, has been further enhanced.
- Now you can create mirror databases and upload parquet and CSV files directly via drag-and-drop, making data replication even easier.

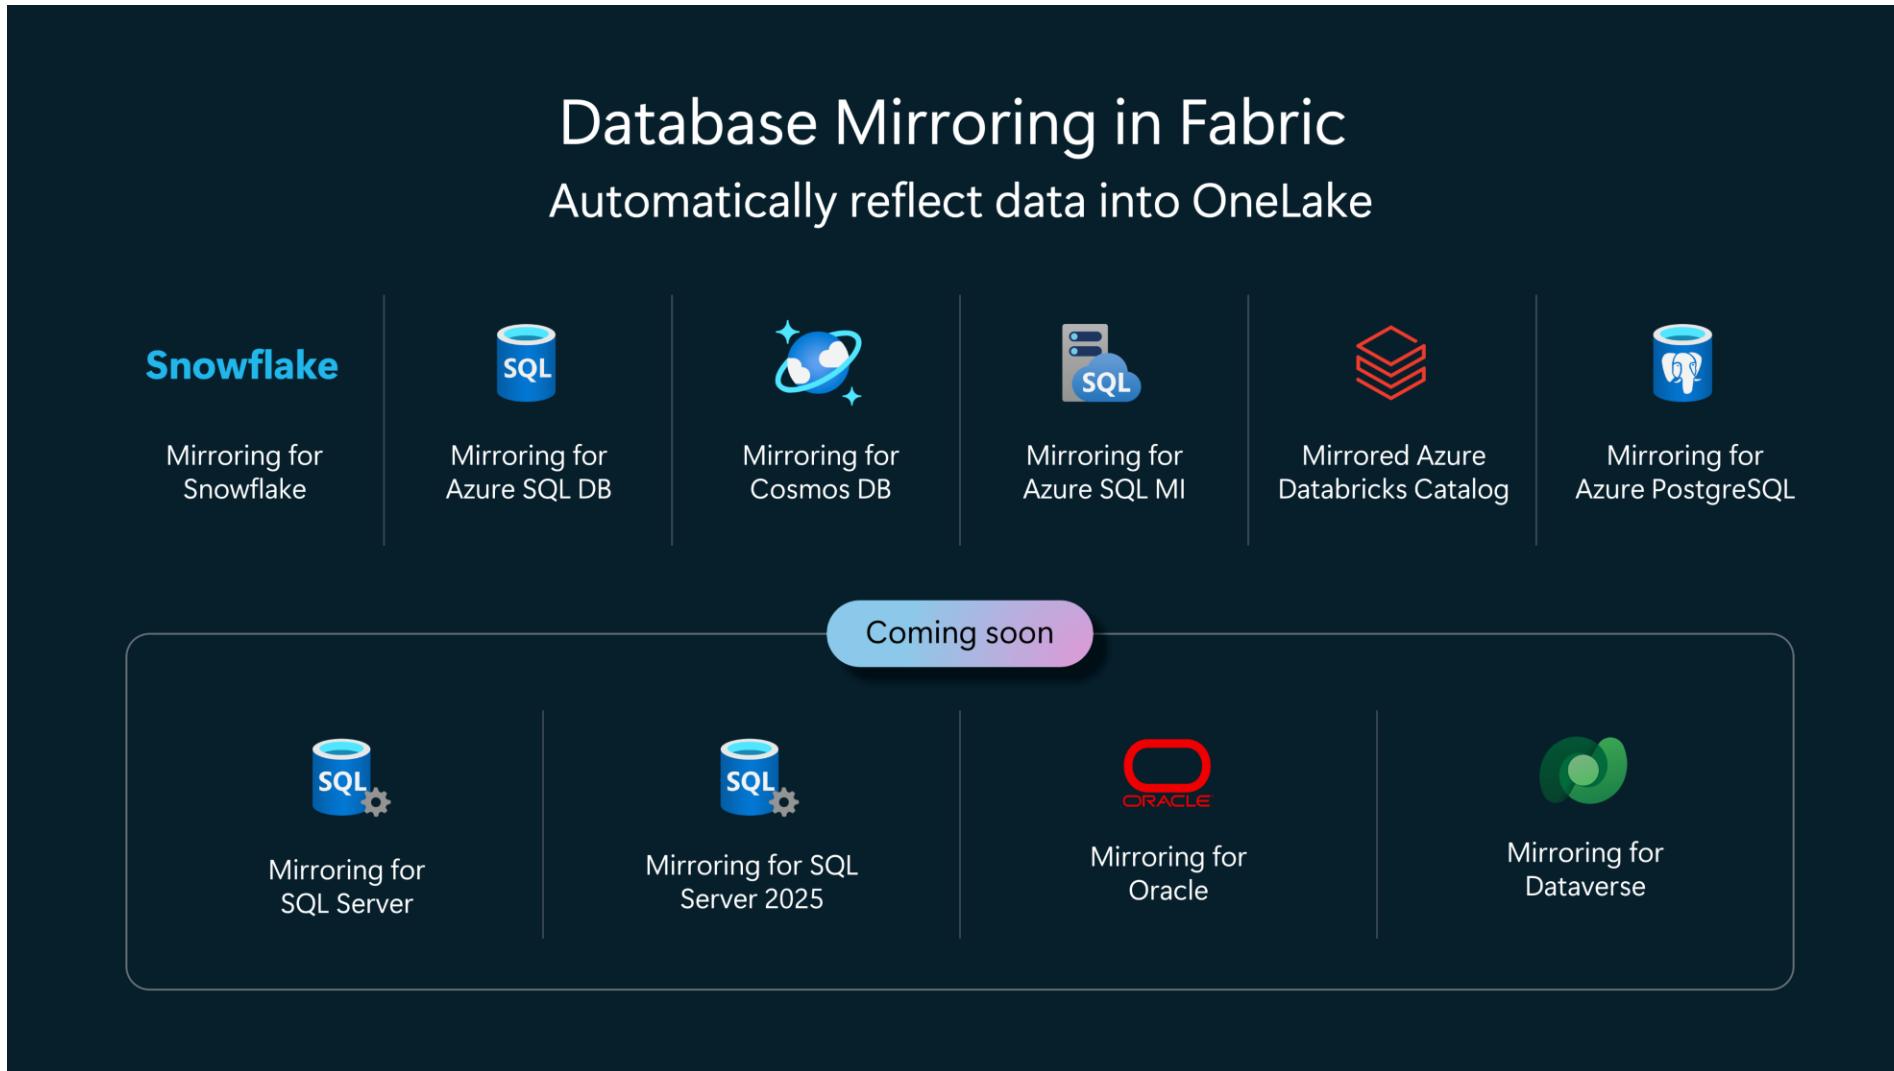
Mirroring for Azure SQL Database: Relevant news

- It now supports the replication of tables without a primary key, providing greater flexibility.
- It reduces the SQL roles and permissions required to enable mirroring, simplifying configuration and increasing security.

Free storage and compute for replicas

- Microsoft offers free storage space for replicas proportional to the capacities purchased (e.g.: 1 TB free for each unit of capacity purchased).
- The compute required to replicate data to OneLake is free and does not consume capacity, while the compute for subsequent data queries (SQL, Spark, Power BI) is charged normally.

Roadmap



Link

- [Mirroring - Microsoft Fabric | Microsoft Learn](#)
- [Open Mirroring \(Preview\) - Microsoft Fabric | Microsoft Learn](#)
- [Announcing Mirroring for Azure Database for PostgreSQL in Microsoft Fabric for Public Preview | Microsoft Community Hub](#)

Unlock the power of Real-Time
Intelligence in the Era of AI: why
Fabric Real-Time Intelligence is a
game-changer

Intro

- Since Microsoft Fabric Real-Time Intelligence was made available in November 2024, it has been adopted by more than 8,900 customers in just a few months
- Companies today recognize the strategic importance of real-time data to respond quickly to change and gain real competitive advantages

Why is this important?

- **Immediate competitive advantage**

Faster decisions, process optimization and timely responses to customer needs

- **Built-in AI**

Integration with real-time data allows you to anticipate problems, detect anomalies, and identify opportunities before it's too late

- **Applicable to all industries**

Not only finance or manufacturing, but also retail, healthcare, transportation, call centers and many other sectors benefit from the ability to act quickly based on timely and accurate data

The main barriers to the past

Historically, companies have found it difficult to implement real-time solutions due to:

- Fragmented and complex data
- Legacy systems using batch processing (delayed data processing)
- High costs and difficulties of infrastructure management

Microsoft Fabric overcomes these limitations with a simple, unified, and integrated solution

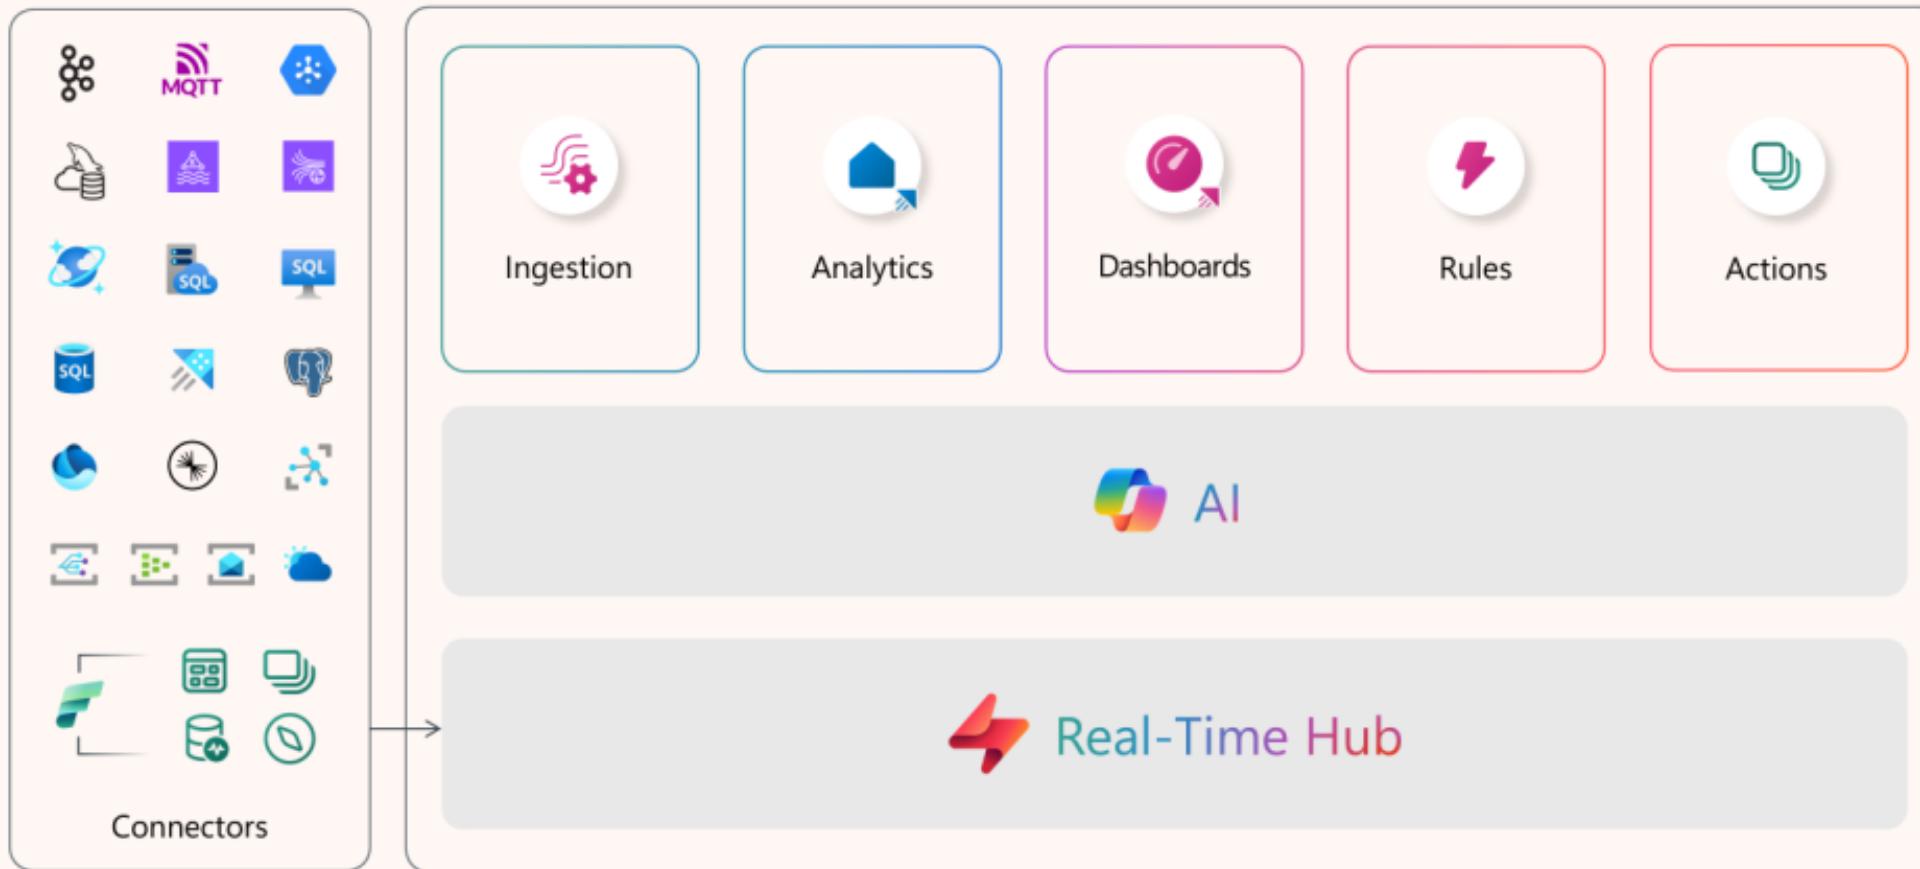
What exactly is RTI in Microsoft Fabric?

It is a complete platform that allows you to:

- **Ingest real-time data** through **Eventstreams**
- **Monitor conditions and react immediately** with **Activator**
- **Analyze and visualize data** using **Eventhouse**, **KQL querysets**, and **Real-Time Dashboards** without the need to know complex languages
- **Manage streaming data sources** through the **Real-Time Hub**, with direct access to OneLake
- Integrate an **AI layer** to generate insights simply by interacting in natural language



Real-Time Intelligence in Microsoft Fabric



Planet-scale infrastructure for real-time data

10 EB

Events and logs per month

350 PB

Ingested daily

5.1 B

Real-time queries per day

19.2 T

Streaming events processed monthly

99.9992%

Success rate across messaging



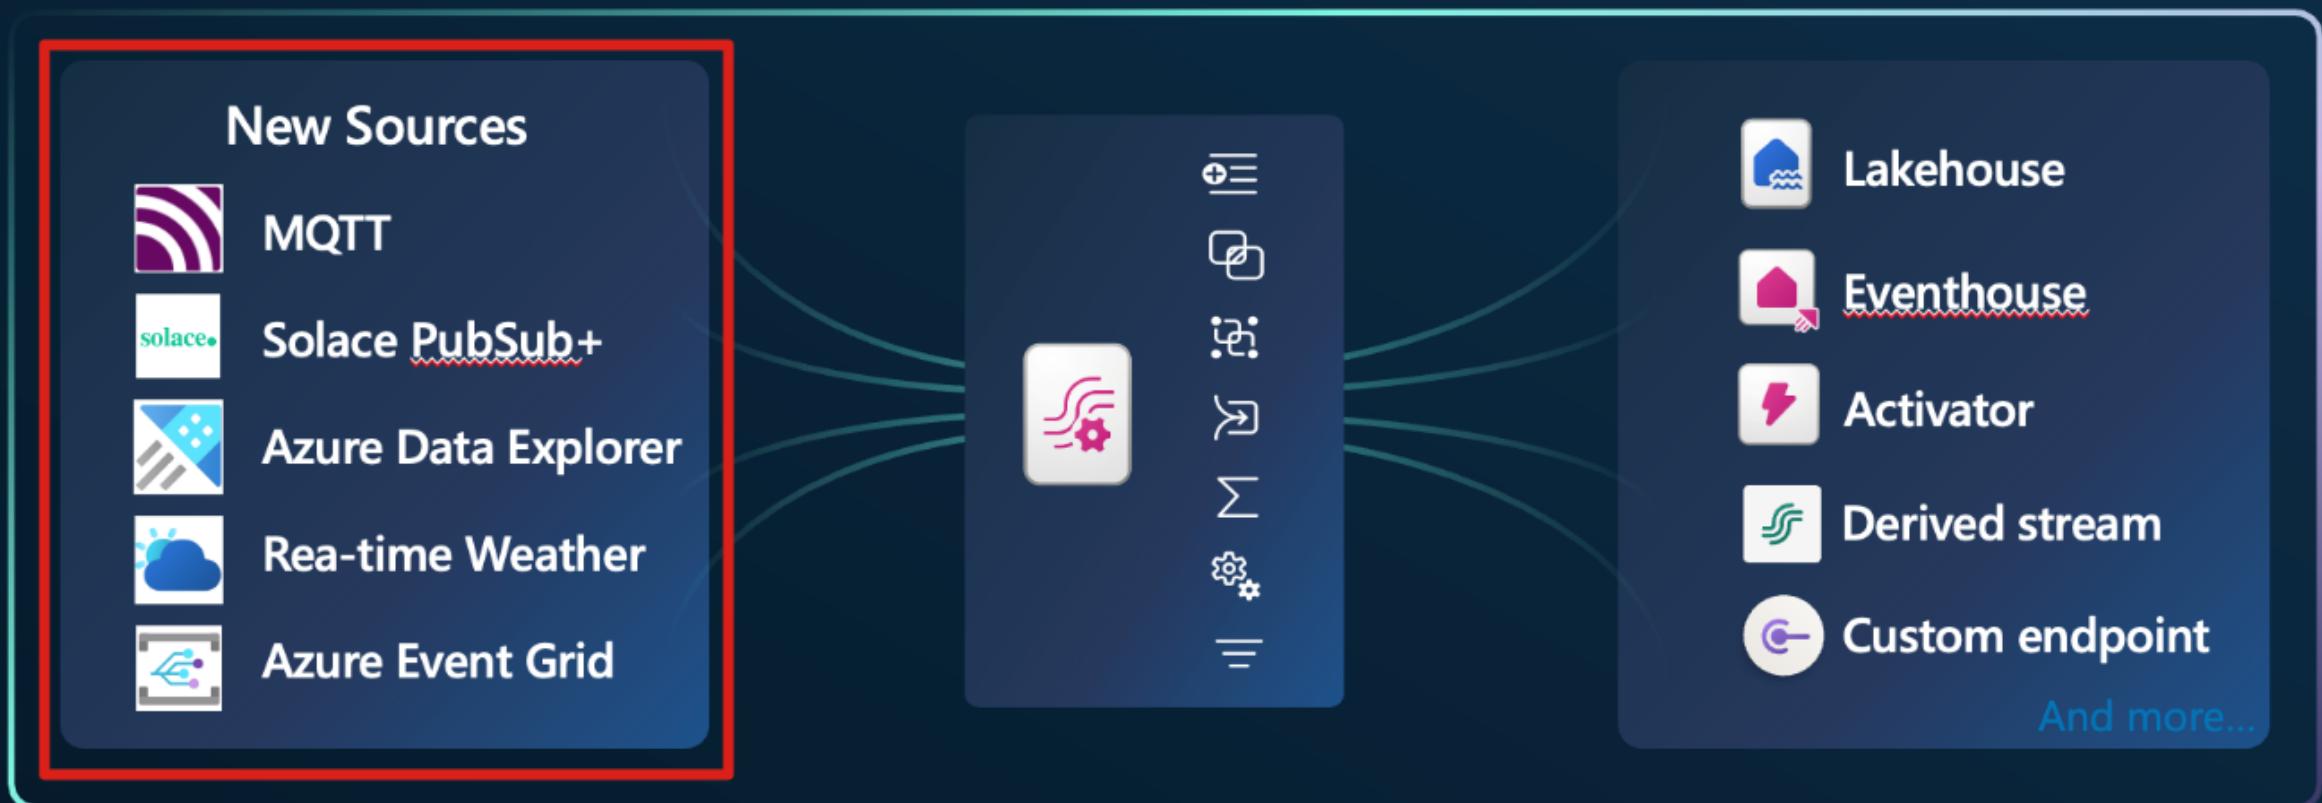
Main innovations presented at
FabCon Vegas 2025

What's new for developers (Developer Ready)

- **New streaming connectors** in preview for **Eventstreams**:
 - MQTT, Solace PubSub+, Azure Data Explorer, Real-time Weather, Azure Event Grid, dati finanziari (es. S&P500).
- **Multiple schema support** to handle different data sources and transformations.
- **Fabric Events (GA)**: Fabric-generated event management (OneLake, Azure Blob, Job events) with advanced filtering capabilities.
- **CI/CD (GA)**: Full support for Continuous Integration/Continuous Deployment with Git, improving collaboration and automation in the development of streaming solutions.



Real-Time Intelligence: Eventstream



Introducing 5 new Eventstream sources for seamless data ingestion and real-time transformation across various data streams

Security news (Secure)

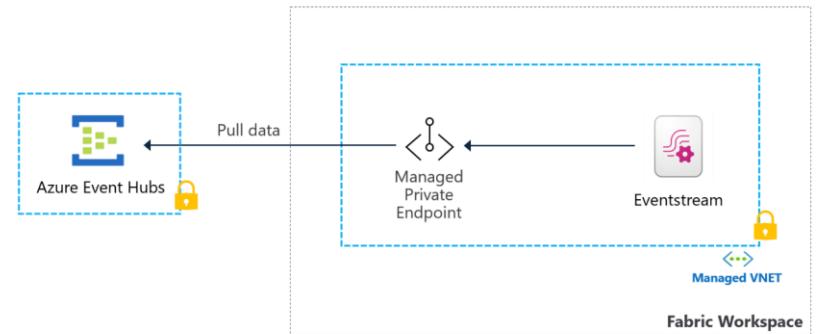
- **Managed Private Endpoint (MPE) in anteprima:**

- Enables secure, private connections to firewalled resources directly from Fabric Eventstream.

- **Microsoft Login ID Authentication in Preview:**

- Increase security and simplify authentication by eliminating the use of SAS keys. Allows you to use managed identities for Azure services such as Logic Apps.

Eventstream pulls data from Azure Event Hub using Managed Private Endpoints.



General improvements

- **Easily manage Activator alerts in Power BI**
 - Directly within reports, without changing interfaces
- **KQL entity diagrams**
 - Graphical visualization of relationships between entities such as tables, functions, and policies

Key benefits for businesses

- **Easy to implement**
 - No more complex or expensive architectures
- **Cost reduction and performance improvement**
 - Real-world cases dramatically reduced implementation costs and time with Fabric

Link

- [Explore Fabric events in Fabric Real-Time hub - Microsoft Fabric | Microsoft Learn](#)
- [What's new and planned for Real-Time Intelligence in Microsoft Fabric - Microsoft Fabric | Microsoft Learn](#)

Fabric Data Factory Summary and Main News



Data Factory in Microsoft Fabric



Best-in-class
Connectivity

Data Movement

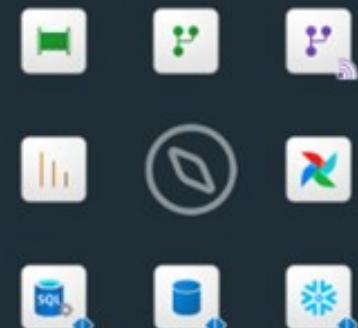
Copy job Pipeline

Orchestration

Pipeline Apache Airflow job

Transformation

Dataflow Gen2 Notebook



Deployment &
Observability



OneLake

AI-powered Intelligence

Security and integration with Azure Key Vault

- Preview support for using secrets stored in **Azure Key Vault** for data connections.
- You can now securely manage your connection credentials using AKV.

Dataflows and advanced orchestration

- Support for **parameters** in **Dataflows** and Dataflow activities within pipelines, making it easier to manage metadata-driven data solutions.

CI/CD and DataOps enhancements

- **General availability (GA)** of CI/CD support in Data Factory:
 - New variable libraries that make it easier to change values between environments.
 - Authentication via Service Principal (SPN) for CRUD API pipelines.

Enterprise-class connectivity and data movement

- Connectivity to more than **170 data sources**.
- Improved support for moving data between **virtual networks (VNets), on-premises, and multi-clouds**.
- Il Virtual Network Data Gateway ora supporta **Data Pipelines, Copy Jobs e Mirroring**.

Copy Job (GA) News

- New data sources (over 20 new connectors).
- Public API and CI/CD support.
- Support for **upsert operations** to SQL databases and **overwrite** operations to Lakehouse tables.
- Real-time monitoring of the copy process.

Advanced mirroring to OneLake

- Continuous replication is now also supported for databases behind firewalls and on-premises (**Azure SQL DB**, soon to be **Snowflake**, and **Azure SQL Managed Instance**).
- Important news such as:
 - Mirroring per **Azure Database for PostgreSQL Flexible Server**.
 - CI/CD support, workspace monitoring, source schema replication, Delta Column Mapping support, UI enhancements for Open Mirroring, and CSV support.
 - Support for tables without primary keys and reduced SQL roles for Azure SQL DB.
 - Mirroring service geo-expansion (new regions supported)
 - Tra cui Italy North [?](#)

Dataflows Gen2 (Enterprise Data Transformations and Orchestration)

- General availability for CI/CD in Dataflows.
- Incremental refresh for optimized, faster updates.
- Significant user experience improvements (multitasking UX, simplified saving and scheduled refreshes).

Advanced innovations in Data Pipelines

- New pipeline trigger based on OneLake events.
- Public preview support for **User Defined Functions (UDF)**.
- Private preview for migrating **SSIS** packages directly to Fabric.
- Full support for **Apache Airflow** to create pipelines in Python format (DAG).
- Pipeline activity limit extended to **120 activities per pipeline**.

AI-Powered Development with Copilot (GA)

- General availability of Copilot on Fabric Data Factory.
- Creation and management of Dataflows and Data Pipelines using textual inputs in natural language.

Easier migration paths to Fabric Data Factory

- New guides and tools to easily migrate from:
 - **Azure Data Factory (ADF)** e **Azure Synapse Pipelines** verso Fabric Data Factory.
 - **Dataflow Gen1 (Power BI Dataflows)** to **Dataflow Gen2** with a simple "Save As" function.

Link

- [Migrate from Azure Data Factory to Fabric Data Factory - Microsoft Fabric | Microsoft Learn](#)
- [Bring your Azure Data Factories into your Microsoft Fabric Workspace](#)
- [Migrate from Dataflow Gen1 to Dataflow Gen2 - Microsoft Fabric | Microsoft Learn](#)
- [Dataflow Gen2 with CI/CD and Git integration - Microsoft Fabric | Microsoft Learn](#)

Autoscale Billing per Spark in Microsoft Fabric

Intro

- Microsoft introduces the new Autoscale Billing model for Spark workloads in Microsoft Fabric.
- This mode is designed to provide greater flexibility and optimize costs by allowing you to pay only for what you consume when running Spark jobs, without directly impacting the fabric capacity you purchase (F2 or higher).

Autoscale Billing vs. Capacity-Based Model

- Two complementary options

| Characteristic | Fabric Capacity Model | Autoscale Billing (Spark) |
|------------------|-----------------------|----------------------------------|
| Billing methods | Fixed cost per tier | Pay-as-you-go |
| Resource scaling | Shared resources | Independent scaling for Spark |
| Resource sharing | Possible contention | Spark job resources |
| Ideal scenario | Predictable workloads | Dynamic and intensive Spark jobs |

Key Benefits of Autoscale Billing

- **Cost efficiency**
 - You only pay for the real duration of your Spark jobs.
- **Independent scaling**
 - Spark jobs scale without disrupting other workloads in Fabric.
- **Flexible quota management**
 - Ability to request additional resources through Azure Quota Management.
- **Dedicated resources for Spark**
 - No interference with other shared workloads.

-  Home
-  Create
-  Browse
-  OneLake
-  Apps
-  Metrics
-  Monitor
-  Learn
-  Real-Time
-  Workloads
-  Workspaces
-  My workspace
- ...

| | Name | Type | Opened | Location | Endorsement | Sensitivity |
|--|---------------------------------|-------------|-------------|-------------------------|-------------|-------------------------------|
| | DE_DEV_E_ADU Report | Report | 3 days ago | DS_DE_DEV E Reporung | — | Confidential\Any User (N...) |
| | PM Spark Benchmarks | App | 8 days ago | Apps | — | — |
| | DE DS AI Tracking | Report | 8 days ago | DS_DE_DevE Reporting | — | Confidential\Any User (N...) |
| | My workspace | Workspace | 9 days ago | Workspaces | — | — |
| | PS_UMI-AOS | Report | 10 days ago | — | — | Confidential\Microsoft Ext... |
| | Azure Synapse VHD Releases | App | 10 days ago | Apps | — | — |
| | Notebook 1 | Notebook | a day ago | fabconworkspace | — | Confidential\Microsoft Ext... |
| | Notebook 2 | Notebook | a day ago | fabconworkspace | — | Confidential\Microsoft Ext... |
| | eventstream | Eventstream | a day ago | fabconworkspace | — | Confidential\Microsoft Ext... |
| | DataSamplingNotebook | Notebook | a day ago | fabconworkspace | — | Confidential\Microsoft Ext... |
| | Notebook 1 | Notebook | a day ago | santhoshclosedworkspace | — | Confidential\Microsoft Ext... |
| | lh | Lakehouse | a day ago | santhoshopenworkspace | — | Confidential\Microsoft Ext... |
| | Notebook 1 | Notebook | a day ago | featureautoscale | — | Confidential\Microsoft Ext... |
| | env | Environment | 8 days ago | featureautoscale | — | Confidential\Microsoft Ext... |
| | envwitheventlisteners | Environment | 8 days ago | fabconworkspace | — | Confidential\Microsoft Ext... |
| | Sample_Lakehouse_887 | Lakehouse | 8 days ago | fabconworkspace | — | Confidential\Microsoft Ext... |
| | RunAnalysisLH | Lakehouse | 8 days ago | fabconworkspace | — | Confidential\Microsoft Ext... |
| | Notebook 7 | Notebook | 9 days ago | msbench | — | Confidential\Microsoft Ext... |
| | Notebook 9 | Notebook | 13 days ago | msbench | — | Confidential\Microsoft Ext... |
| | LH | Lakehouse | 18 days ago | fabconworkspace | — | Confidential\Microsoft Ext... |
| | notebook_for_handling_data_skew | Notebook | 21 days ago | fabconworkspace | — | Confidential\Microsoft Ext... |

How Autoscale Billing works

Budget control

- Administrators can set a maximum limit of CUs (Capacity Units) to ensure cost control.

Transparent resource allocation

- Spark workloads no longer consume shared fabric capacity.
- Once the CU limit is exceeded, Spark jobs are queued (batch jobs) or slowed down (interactive queries).
- Usage tracked separately in the Capacity Metrics App.

Spend tracking and optimization

- Total visibility into your spend in Azure Cost Analysis.
- Ability to filter Spark consumption for precise cost control.

New monitoring interface

Autoscale Billing page

- New dedicated page in the Capacity Metrics App for:
 - Monitor Spark usage across workspaces.
 - Analyze consumption trends over time.
 - Investigate job duration and resource allocation.

Quota management for enterprise users

- Enterprise users can request a quota increase through Azure Quota Management, then configure the limits directly from the Fabric Capacity Settings page.

(Available initially in the UK South region, with global rollout by April 3.)

Link

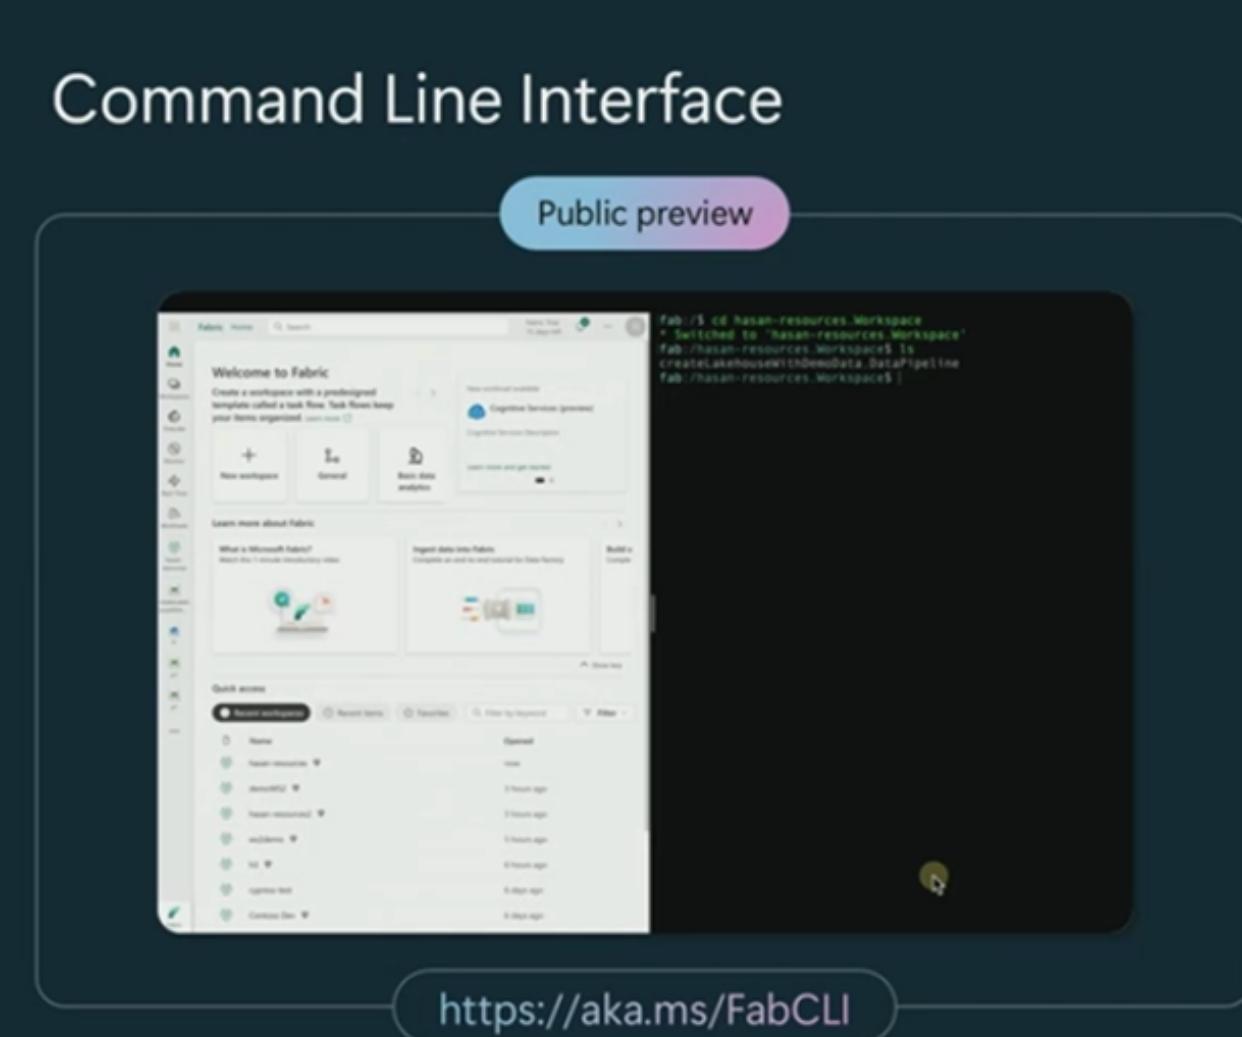
- [Autoscale Billing for Spark in Microsoft Fabric - Microsoft Fabric | Microsoft Learn](#)
- [Microsoft Fabric - Pricing | Microsoft Azure](#)

Professional Dev Platform

Fabric CLI

Command Line Interface

Public preview



The screenshot shows the Microsoft Fabric Command Line Interface (CLI) in a public preview. On the left, there's a browser-based interface with a sidebar for 'New workspace', 'General', 'Basic data analytics', and 'Import data into Fabric'. The main area shows 'Cognitive Services (preview)' and 'Cognitive Service Description'. On the right, a terminal window displays a command-line session:

```
fab:/ $ cd hasan-resources.Workspace
fab:/ $ cd hasan-resources.Workspace
* Switched to "hasan-resources.Workspace"
fab:/hasan-resources.Workspace$ ls
createLakehouseWithDemoData, DataPipeline
fab:/hasan-resources.Workspace$
```

At the bottom, a URL is provided: <https://aka.ms/FabCLI>.

A command-line interface that allows users to interact with Microsoft Fabric in a familiar, file system-like way.

Complete repetitive tasks with CLI commands. No need to jump to the browser!

Use your favorite scripting language to customize CI/CD workflows

Key features

The Fabric CLI allows you to:

- Run Notebooks
- Manage data pipelines
- Administer permissions and security
- Automate tasks and workflows in Fabric

Objective: To provide a viable alternative to graphical user interfaces, similar to what the Azure CLI is for Azure

The new Fabric CLI adopts Windows and Linux command syntax, making it easy for developers to use familiar commands such as ls and cd to navigate and manage workspaces and resources. This tool is especially useful for developers who prefer to work with the keyboard and automate their development processes.

Two main modes of use

Interactive mode

- Users execute individual commands one at a time with immediate response.
- Ideal for exploration, learning, and quick activities.

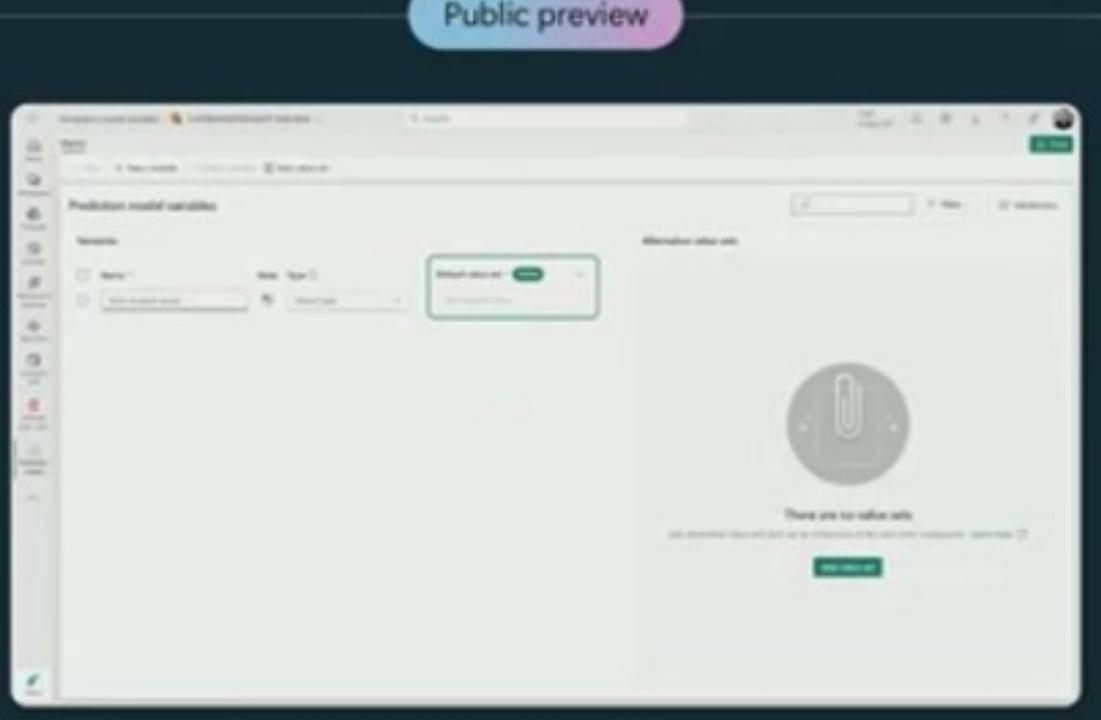
Command line mode (scripting/batch)

- It allows multiple commands to be executed in sequence, which is useful for advanced automation.
- Perfect for scripts, integrations, and CI/CD pipelines.

Variable Library

Variable Libraries

Public preview



The screenshot shows a user interface for managing variables. At the top, there's a header with the title 'Variable Libraries' and a 'Public preview' badge. Below the header is a search bar and a sidebar with various icons. The main area displays a table with columns for 'Name', 'Type', and 'Value'. A prominent green button labeled 'Create variable' is located at the bottom right of this table. In the center of the page is a large circular icon with a paperclip symbol, accompanied by the text 'There are no active vars'.

New Fabric item to define and manage variables at the workspace level

Create environment-specific configurations for different deployment stages

Centralized management with support for multiple variable types

Use in pipelines, notebooks, shortcuts for lakehouse with more integrations coming

Variable Library in Microsoft Fabric (anteprima)

- Designed to simplify and centralize configuration management in Continuous Integration/Continuous Deployment (CI/CD) processes in Fabric, avoiding hard-coded values and facilitating switching between environments (development, test, production)
- VL is a new type of item that allows users to define and manage workspace-level variables, making them available to all items in the workspace, such as data pipelines, notebooks, and Lakehouse Shortcuts

Key features and benefits

- **Environment-specific configurations:**
Define different values for each environment (development, test, production) to facilitate and speed up releases.
- **Centralized management:**
All variables are managed from a central location, making it quick and easy to update and maintain.
- **CI/CD pipeline integration:**
Supports Git integration and deployment pipelines, with APIs available for full automation.
- **Multiple variable types supported:**
Including booleans, integers, numbers, strings, GUIDs, and DateTime for maximum flexibility.

Integration with other Fabric elements

- **Data pipeline:** Use of variables in dynamic content fields.
- **Notebook:** Native use of variables directly in the notebook code.
- **Lakehouse Shortcut:** Dynamic management of connection configurations and paths using variables.

User Data Function

- You can now create functions that can be reused throughout your system, using languages such as Python and C.
- This provides a software as a service (SAS) experience on functions, improving code productivity and reusability.

User data functions

Public preview

```
cursor = sqlconnection.cursor()
cursor.execute("UPDATE [dbo].[Opportunity] SET [Discount] = ?, [Comments] = ? WHERE [Opportunity_Number] = ?;", (discount, comments, opportunitynumber))
sqlconnection.commit()
sqlconnection.close()
logging.info('Opportunity (opportunitynumber) was updated.')
return f'Opportunity (opportunitynumber) was updated.'
```

```
if opportunitystatus == 1 or status == 3:
    logging.info('Opportunity cannot be changed.')
    return f'This opportunity cannot be changed since it is either in Won or Lost state.'
else:
    # Opportunity needs to be updated in the table
    sqlconnection = engine.connect()
    sqlconnection.begin()
    cursor = sqlconnection.cursor()
    cursor.execute("UPDATE [dbo].[Opportunity] SET [Discount] = ?, [Comments] = ?, [Opportunity_Number] = ?, [Discount] = ?, [Comments] = ? WHERE [Opportunity_Number] = ?;", (discount, comments, opportunitynumber, discount, comments, opportunitynumber))
    sqlconnection.commit()
    sqlconnection.close()
    logging.info('Opportunity (opportunitynumber) was updated.')
    return f'Opportunity (opportunitynumber) was updated.'
```

Write and run reusable custom business logic integrated with your resources Fabric

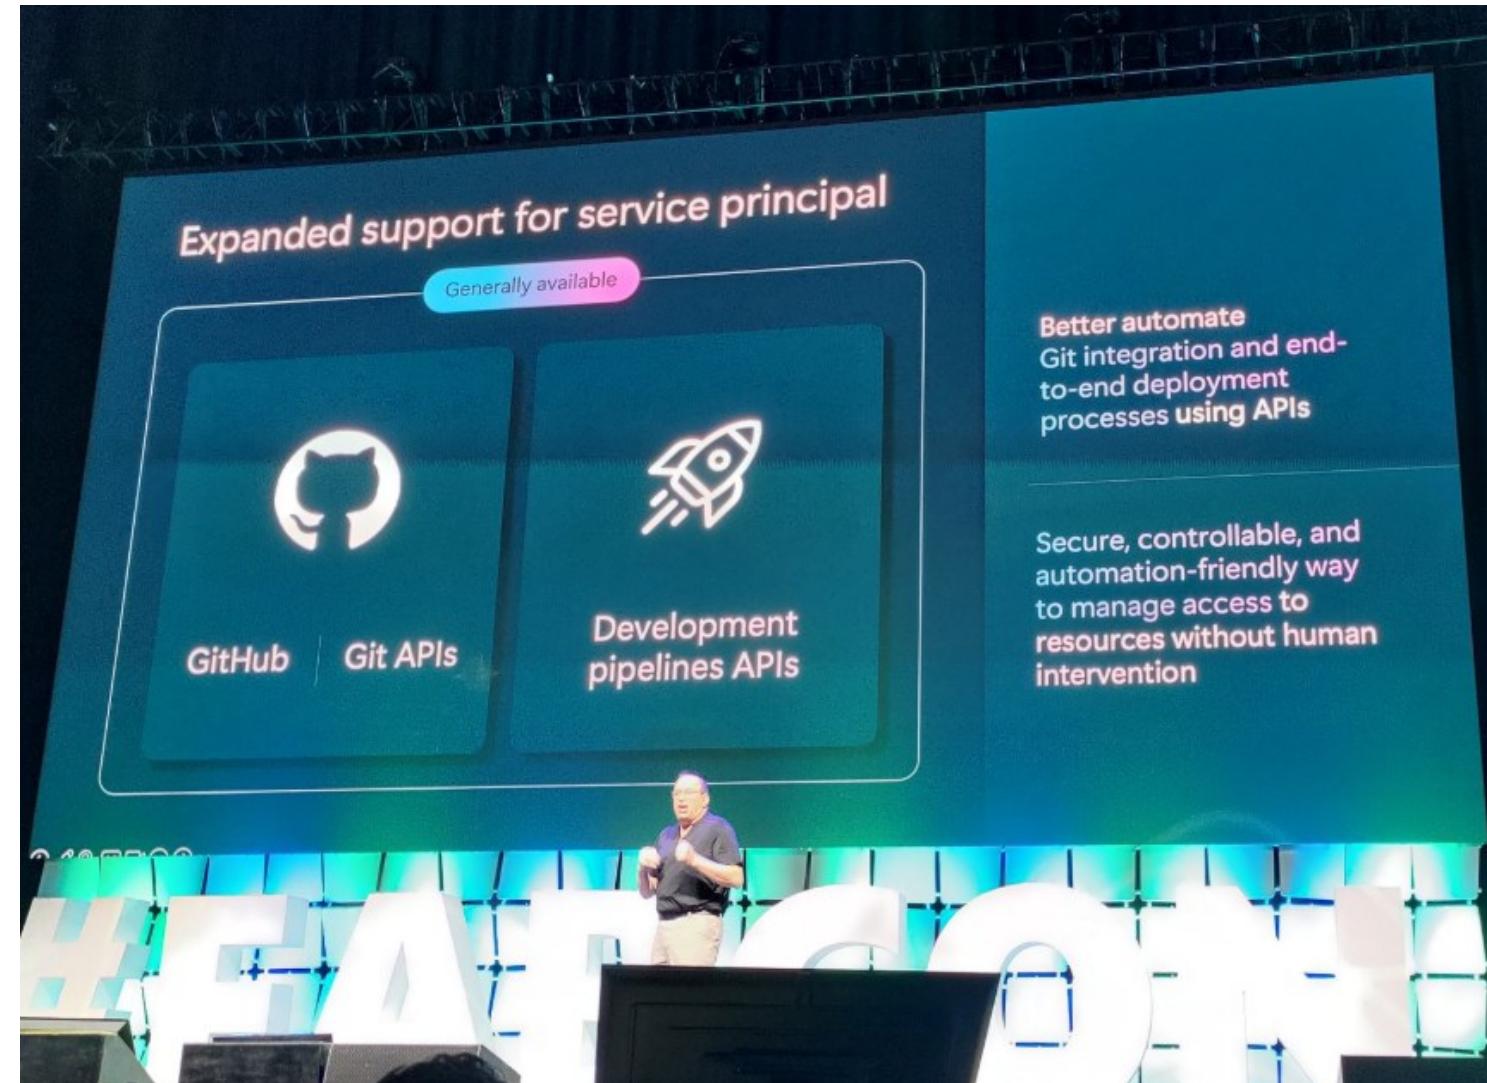
Streamline data engineering, data validation, cleansing, and integrations with external systems

Edit, debug, and publish your data functions using VS Code

FABRIC

Expanded support for Service Principal

- Service Principal Name (SPN) support: Developers can now use SPN identities when working with GitHub and GitHub APIs, when developing pipelines, and using APIs.
- This support has been extended to all platform scenarios.



Link

- [fabric-cli | 🔥 Microsoft Fabric CLI](#)
- [Fabric March 2025 Feature Summary | Microsoft Fabric Blog | Microsoft Fabric](#)

Power BI

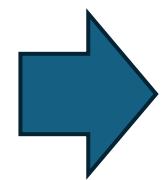
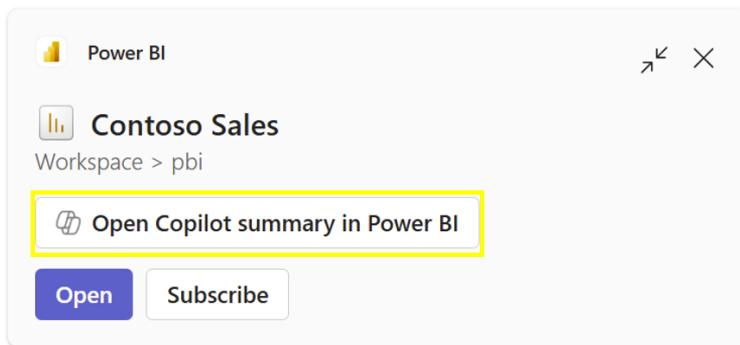
Copilot auto-summary directly from Teams and Outlook

- You can now generate an automatic Copilot summary of a Power BI item (such as a report or dashboard) directly from a Teams or Outlook message, with a single click on the "Open Copilot summary in Power BI" button.
- The button will only be visible to users with Copilot enabled and to items that support Copilot.
- The feature is available on all Teams and Outlook platforms (desktop, browser, mobile).

Key benefits:

- Quick access to summary reports and dashboards.
- Easier to understand at a glance the data shared.

Copilot auto-summary directly from Teams and Outlook



The screenshot shows the Microsoft Copilot summary in Power BI. The main area displays various dashboards and reports for "Contoso Sales". On the right side, there is a sidebar titled "Copilot" with the following sections:

- Give me an executive summary
- Anticipate likely questions from leadership
- What insights should I look into?

Below the sidebar, there is a "Contoso Sales Executive Summary" section and a "Sales Performance" section with bullet points:

- The total Revenue Won amounts to \$26.44M with a Win/Loss percentage of 49.6%. Additionally, the Forecast based on the Win/Loss Ratio is \$13.81M.
- The closing percentage by product categories is highest for Warranties.

At the bottom, there is a "Ask a question about this report" input field.

Ad hoc calculation support

- Copilot can now automatically create ad-hoc calculations (measures and calculated columns) to answer specific questions about the data, even if these measures do not exist in the original model.
- Questions may include:
 - Year-on-year growth in sales.
 - Number of employees hired before a date.
 - Calculation of the ratio between product categories.
 - Customers who have not purchased products.
- Copilot automatically generates DAX queries to answer them accurately.
- Currently only available in the report's Edit mode, not in Read mode.

Ad hoc calculation support

- What was the year-over-year growth for sales?
- How many employees were hired before 2020?
- Calculate the ratio of cosmetic product orders to all products.
- Which customers did not buy any products?

Change in Average Unit Price from Feb 2012 to Mar 2012

Results DAX query

This answer required more analysis than usual. Check it carefully. AI can make mistakes. [Learn more](#)

| [Average Unit Price Feb 2012] | [Average Unit Price Mar 2012] | [Difference] |
|-------------------------------|-------------------------------|--------------|
| 1949.98 | 1761.71 | -188.27 |

Based on the available data, the average unit price decreased from February 2012 to March 2012. The average unit price in February 2012 was 1949.9776, while in March 2012 it was 1761.7123. This represents a decrease of 188.2653.

[Copy table](#)

Change in Average Unit Price from Feb 2012 to Mar 2012

Results DAX query

This answer required more analysis than usual. Check it carefully. AI can make mistakes. [Learn more](#)

```
1 // DAX query generated by Fabric Copilot with "how has the average unit price changed from Feb 2012 to March 2012? Show the difference"
2 DEFINE
3 // Calculate the average unit price for February 2012
4 VAR AvgUnitPriceFeb2012 = CALCULATE(
5     AVERAGE('Internet Sales'[Unit Price]),
6     'Date'[Year] == 2012,
7     'Date'[Month Number Of Year] == 2
8 )
9 // Calculate the average unit price for March 2012
10 VAR AvgUnitPriceMar2012 = CALCULATE(
11     AVERAGE('Internet Sales'[Unit Price]),
12     'Date'[Year] == 2012,
```

This query calculates the average unit price of products sold online for two specific months in 2012: February and March. It then determines the difference between these two averages. First, it calculates the average unit price for February 2012 by considering only the sales data from that month and year. Similarly, it calculates the average unit price for March 2012 using the sales data from that month and year. Finally, it computes the difference between the average unit price in March 2012 and the average unit price in February 2012. The result is a table showing the average unit price for February 2012, the average unit price for March 2012, and the difference between these two averages. This allows you to see how the average unit price

[Run query](#) [Copy query](#)

how has the average unit price changed from Feb 2012 to March 2012? Show the difference

i This answer required more analysis than usual. Check it carefully. AI can make mistakes. [Learn more](#)

Advanced: Check the DAX

Based on the available data, the average unit price decreased from February 2012 to March 2012. The average unit price in February 2012 was 1949.9776, while in March 2012 it was 1761.7123. This represents a decrease of 188.2653.

| [Average Unit Price Feb 2012] | [Average Unit Price Mar 2012] | [Difference] |
|-------------------------------|-------------------------------|--------------|
| 1949.98 | 1761.71 | -188.27 |

[Expand answer](#)

↙ How Copilot arrived at this

Is this response helpful?



Data point annotations in PowerPoint presentations (preview)

- You can now easily add **text annotations to data points** from Power BI charts inserted into PowerPoint presentations.
- The notes:
 - They automatically update with data.
 - They are saved only in the specific instance of the Power BI add-in in PowerPoint, making it possible to annotate different for the same visual in different contexts.
 - They disappear automatically if the annotated data is no longer visible (for filters or data updates).
- Annotations also appear in graph snapshots if they are visible at the time of capture.

Data point annotations in PowerPoint presentations (preview)

The screenshot shows a Microsoft PowerPoint slide titled "Pipeline by Stage" with the subtitle "Revenue Opportunities". The chart is a horizontal bar chart titled "Pipeline by Stage" with the y-axis labeled "Sales Stage" and the x-axis labeled "Opportunity Count" ranging from 0 to 200. The data points are:

| Sales Stage | Opportunity Count |
|-------------|-------------------|
| Lead | ~220 |
| Quality | ~100 |
| Solution | ~70 |
| Proposal | ~40 |
| Finalize | ~10 |

Annotations are present on the chart:

- A callout bubble points to the "Quality" stage bar with the text: "Contoso renewal is now in the Quality stage".
- A callout bubble points to the "Finalize" stage bar with the text: "Closing 14 deals in the last week by the APAC team".

At the bottom left of the slide, there is a note placeholder: "Click to add notes".

At the bottom of the screen, the PowerPoint ribbon is visible with tabs like File, Home, Insert, Draw, Design, etc., and a status bar showing "Slide 2 of 2" and "English (United States)".

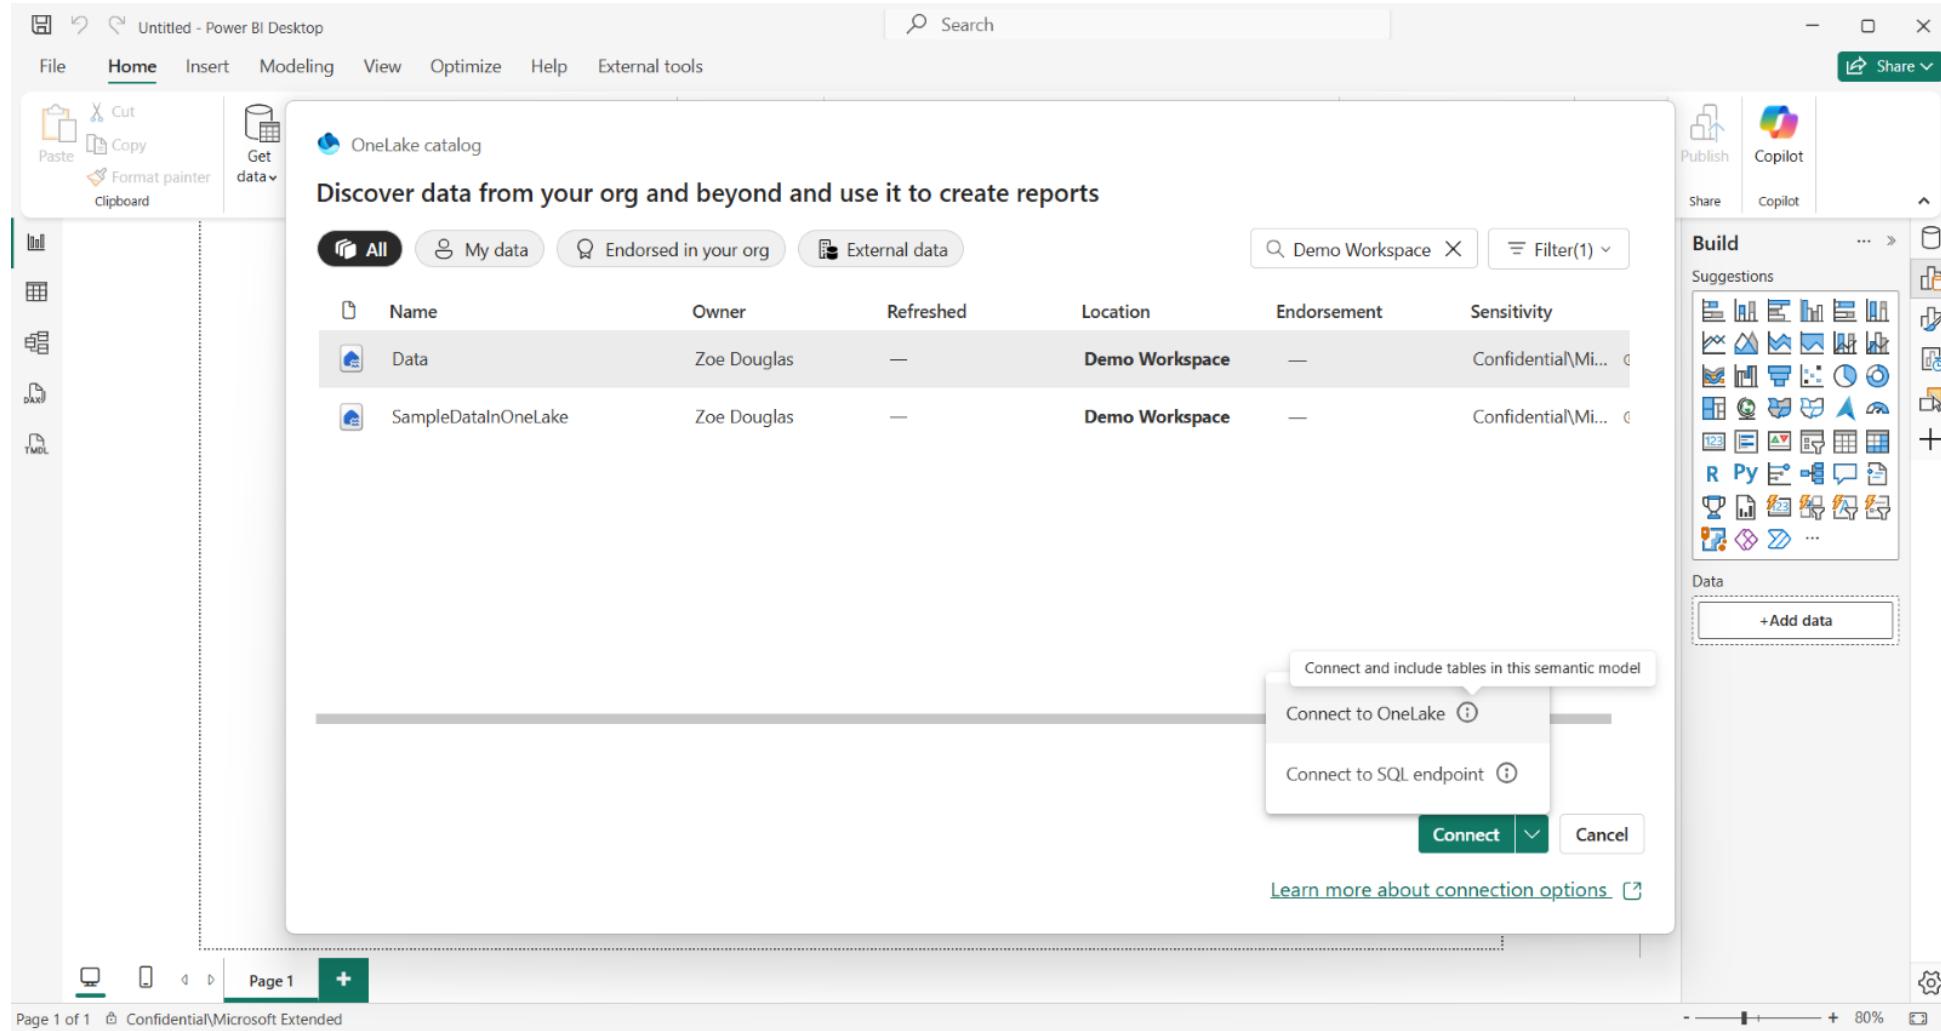
Creating semantic models with Direct Lake in Power BI Desktop

- Announced public preview to create **semantic models** in **Direct Lake** mode directly in Power BI Desktop.
- Direct Lake enables lightning-fast queries on large volumes of data without having to deal with data refresh or duplication.
- For the first time, you can combine tables from multiple OneLake sources (Lakehouse and Warehouse) into the same semantic model.

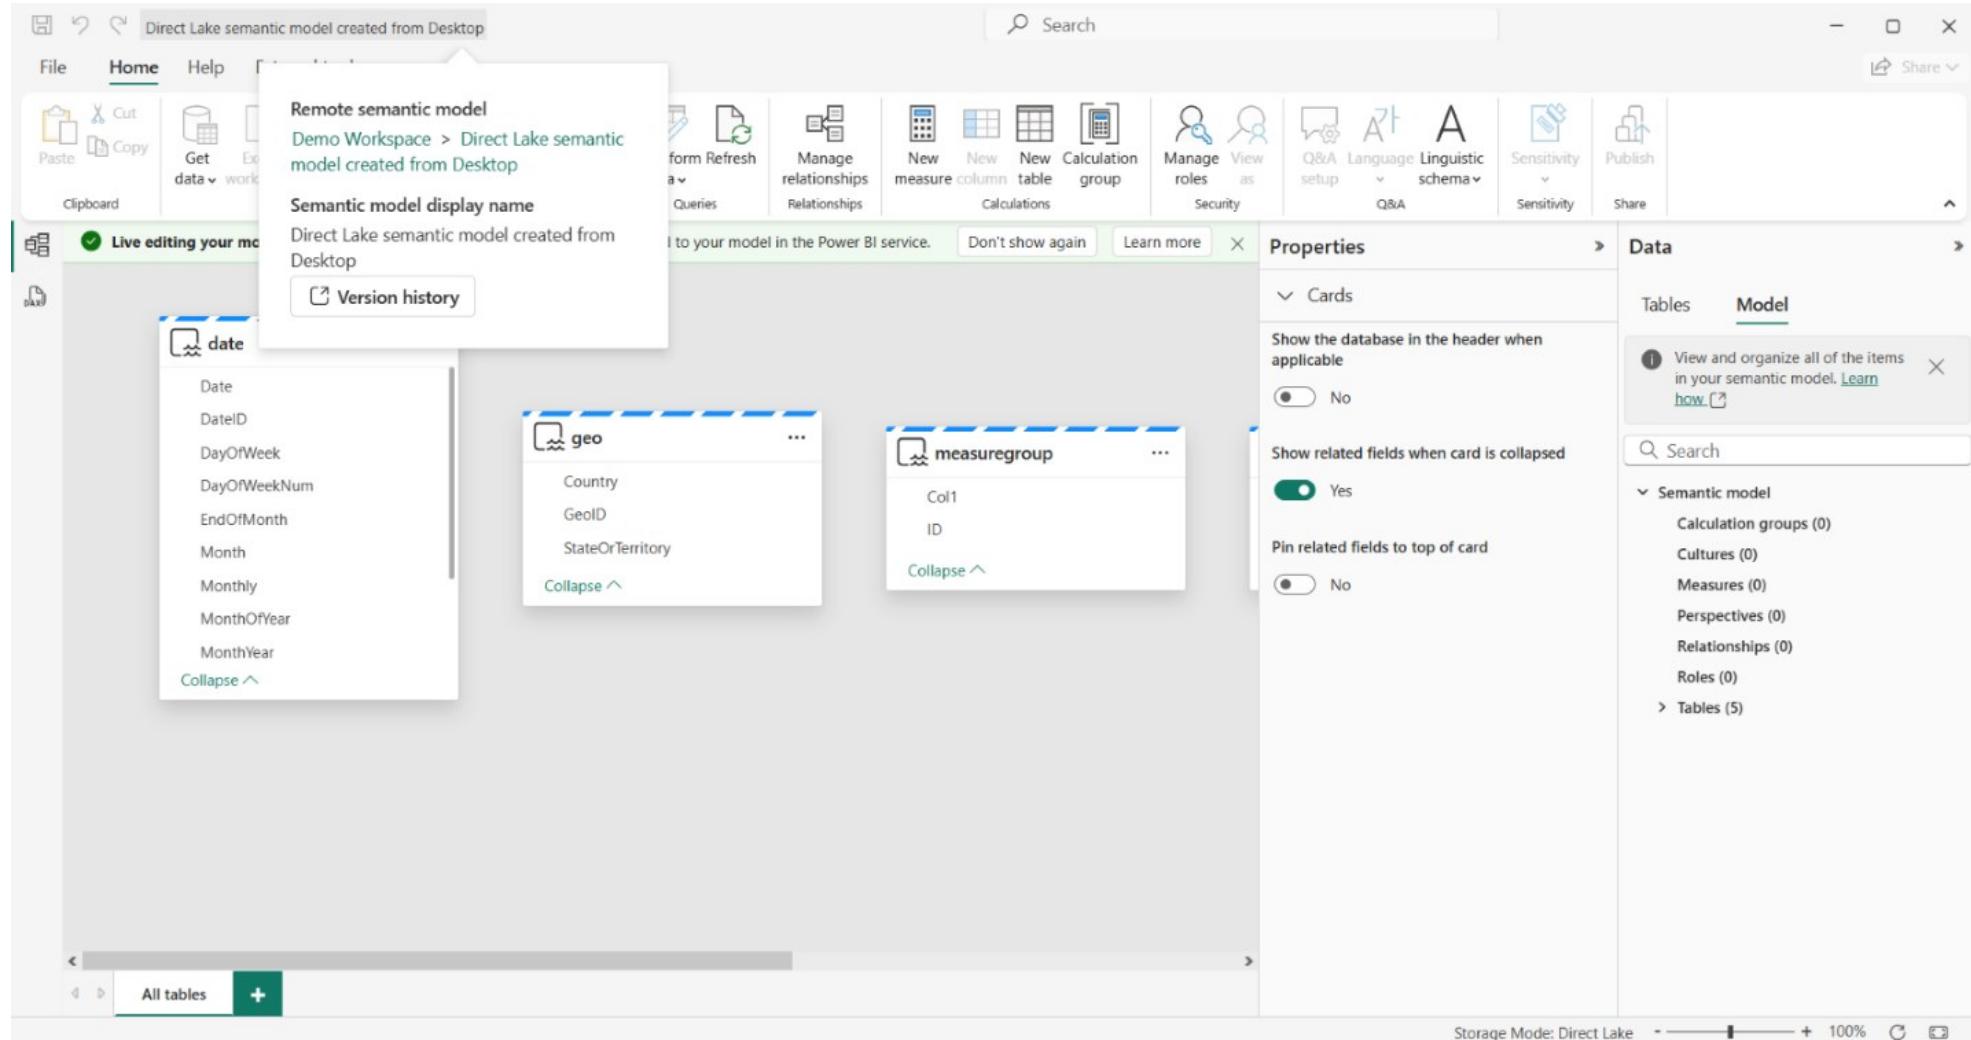
Simple procedure

1. Turn on the preview feature in Power BI Desktop.
2. Select Lakehouse or Warehouse from the OneLake catalog and click "Connect".
3. Create the semantic model by selecting workspaces and tables.
4. Live modify the semantic model by adding additional tables, relationships, measures, hierarchies, and calculation groups.
5. DAX query view is available for quick testing and analysis.

Creating semantic models with Direct Lake in Power BI Desktop



Creating semantic models with Direct Lake in Power BI Desktop



Link

- [Power BI March 2025 Feature Summary | Microsoft Power BI Blog | Microsoft Power BI](#)
- [Create data point annotations \(preview\) - Power BI | Microsoft Learn](#)

New innovations in Microsoft
Purview for protected, AI-ready
data

Enhancing Microsoft Purview Data Loss Prevention (DLP) support for lakehouse in Fabric

- Microsoft Purview enhances its data security capabilities in Fabric, allowing you to block access to sensitive data contained in your Lakehouses to external users ("guests").
- Security administrators will be able to set policies that restrict access to only internal users or data owners when sensitive data, such as personally identifiable information (PII), is detected.
- This significantly reduces the risk of sensitive information leaking to external users.

Enhancing Microsoft Purview Data Loss Prevention (DLP) support for lakehouse in Fabric

The screenshot shows the Microsoft Purview Data Loss Prevention (DLP) interface within the OneLake catalog. The top navigation bar includes the Power BI logo, the catalog name "OneLake catalog", a search bar, and a "Fabric Trial: 20 days left" notification. On the left, a sidebar provides navigation options like "Explore", "Govern (preview)", "All items", "My items", "Endorsed items", "Favorites", "Workspaces", and "Customer_Information". The main content area displays a table with columns "Name", "Type", and "Owner". A single item, "Customer_Info", is listed as a "Lakehouse" type owned by "AdminUser01". To the right, a modal window titled "Sensitive info found" for the item "Customer_Info" is open. It states: "We automatically detected policy issues with this data on 3/16/25. Learn more about data protection policies." Below this, it says: "To find sensitive info in your data, open this lakehouse. Spotted a mistake? Report an issue to the admin, or override the policy label." A red warning icon indicates "Access restricted" with the message: "Restricted to your organization only. This item contains PII data and must not be shared outside the organization. Issues found: Credit Card Number." Buttons for "Report an issue" and "Override" are at the bottom of the modal.

| Name | Type | Owner |
|---------------|-----------|-------------|
| Customer_Info | Lakehouse | AdminUser01 |

Sensitive info found
Customer_Info

We automatically detected policy issues with this data on 3/16/25.
Learn more about data protection policies

To find sensitive info in your data, open this lakehouse. Spotted a mistake?
Report an issue to the admin, or override the policy label.

Access restricted

Restricted to your organization only
This item contains PII data and must not be shared outside the organization.
Issues found: Credit Card Number.

Report an issue Override

Expanding DLP policy support for additional Fabric items

- Purview DLP now supports user notifications (via "policy tips") not only for Lakehouse, but also for KQL databases (used for real-time analytics) and replicated (mirrored) databases from sources such as Azure Cosmos DB, Azure SQL, Azure Databricks, and Snowflake.
- The extension of this feature helps users easily identify when they are working with sensitive data, reducing the risk of accidental loss and strengthening security during data transfers within the Fabric environment.

Expanding DLP policy support for additional Fabric items

The screenshot shows the Power BI OneLake catalog interface. On the left, there's a sidebar with various navigation options like Home, Copilot, Create, Browse, OneLake, Apps, Metrics, Monitor, Workspaces, Customer Information, Azure Mirrored, and a Power BI icon at the bottom. The main area is titled "OneLake catalog" and shows a list of items under "Explore". A dropdown menu indicates the domain is "Customer Information". The list includes:

| Name | Type | Owner |
|----------------------|--------------------------|-------------|
| Azure-Mirrored | Mirrored database | AdminUser01 |
| Azure-Mirrored | Semantic model (default) | AdminUser01 |
| Customer_Info | Semantic model (default) | AdminUser01 |
| Customer_Information | Eventhouse | AdminUser01 |
| Flight Customer Info | Semantic model | AdminUser01 |
| Customer_Information | KQL Database | AdminUser01 |

To the right, a modal window titled "Sensitive info found" for "Customer_Information" is open. It states: "We automatically detected policy issues with this data on 3/11/25. Learn more about data protection policies." Below this, it says: "To find sensitive info in your data, open this KQL database. Spotted a mistake? Report an issue to the admin, or override the policy label." Under "Policies", it says: "Your org identified these policy issues in the data." A callout box highlights: "This item contains PII data and must not be shared outside the organization." It also lists "Issues found: Credit Card Number." and buttons for "Report an issue" and "Override".

The screenshot shows the Power BI OneLake catalog interface. The left sidebar includes links for Home, Copilot, Create, Browse, OneLake, Apps, Metrics, Monitor, Workspaces, Customer_Information, Azure-Mirrored, and Power BI. The main area displays the 'OneLake catalog' under the 'Explore' tab, with a search bar and a 'Domain: Customer Information' dropdown. A table lists items, with one item ('Azure-Mirrored') highlighted. A tooltip for this item provides details about Purview DLP policy issues.

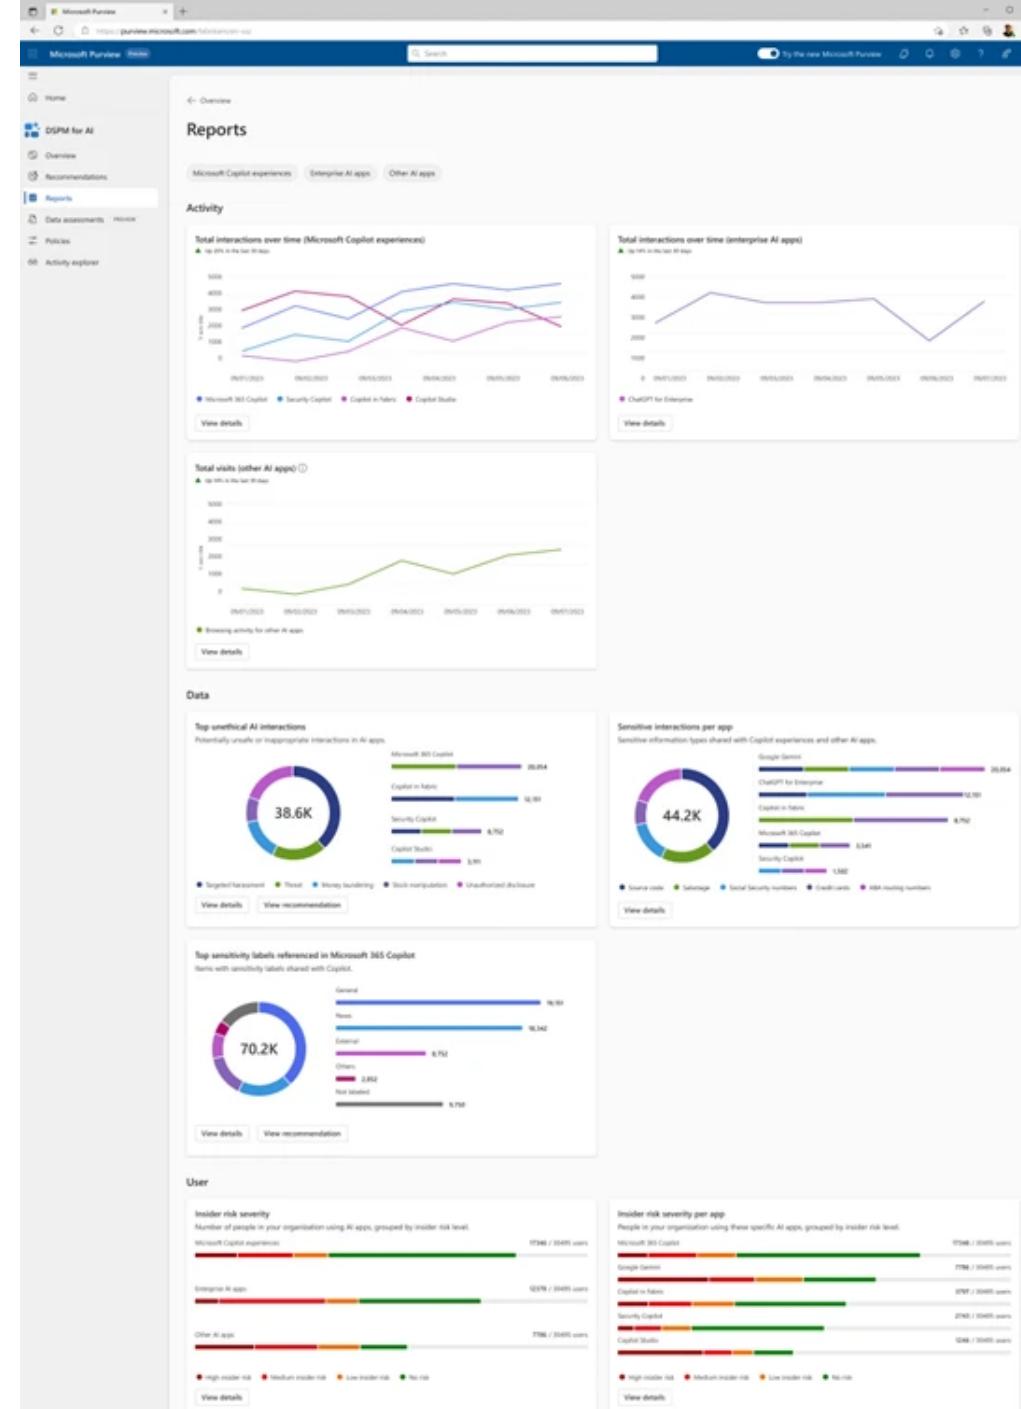
| Name | Type | Owner | Refreshed | Location |
|----------------------|----------------|-------|---------------------|----------------------|
| Azure-Mirrored | Azure-Mirrored | | — | Customer_Information |
| Azure-Mirrored | | | 3/12/25, 4:40:01 PM | Customer_Information |
| Customer_Info | | | 3/10/25, 3:45:37 PM | Customer_Information |
| Customer_Information | | | — | Customer_Information |
| Flight Customer Info | | | 3/12/25, 3:02:49 PM | Customer_Information |
| Customer_Information | | | — | Customer_Information |

- Policy tip triggered by Purview DLP due to PII being discovered in Mirrored databases.

Microsoft Purview for Copilot in Fabric

- Microsoft Purview is integrated into Copilot, starting with Copilot for Power BI, to provide advanced security controls and risk management related to the use of artificial intelligence.
- Key benefits:
 - Identify risks associated with sensitive data ingested into interactions with Copilot, with suggestions for corrective actions via dedicated dashboards (DSPMs).
 - Detection of risky uses of AI (e.g., users accidentally sharing sensitive data or employees using AI to improperly extract sensitive data).
 - Comprehensive monitoring of AI activities through audit tools, legal request management (eDiscovery), data retention policies, and detection of non-compliant behavior.

Microsoft Purview for Copilot in Fabric



Data observability, now in preview, within Microsoft Purview Unified Catalog

- Microsoft Purview Unified Catalog introduces data observability (in preview), allowing users to quickly identify the causes of data quality issues.
- It provides a clear visual representation of the relationships between data assets, governance domains, glossary terms, and data products through a "lineage" view.
- These capabilities allow users to quickly improve data quality, reducing the time it takes to resolve issues and respond more easily to regulatory requirements.

Data observability, now in preview, within Microsoft Purview Unified Catalog

The screenshot displays the Microsoft Purview Unified Catalog interface. On the left, the navigation bar includes Home, Solutions, Learn, Settings, Unified Catalog (which is selected and highlighted in blue), Data Map, Data Lifecycle Management, Data policy, and Data Loss Prevention. The main workspace shows a data lineage diagram. A central node labeled "Organization Data" (Corporate Functions) has two outgoing arrows: one to a node labeled "Opportunities conversion to sales" (Sales), and another to a third node. The third node is partially visible on the right. The "Opportunities conversion to sales" node also has an arrow pointing to the third node. At the top of the workspace, there are buttons for Current view, Data Estate, Hierarchy filter: Data Products, Reset view, Undo, Redo, Scope to selection, Show Lineage, Show non-governed Assets, Save screenshot, and a search bar. To the right of the lineage diagram is a detailed card for "Organization Data". The card includes sections for Status (Draft, Master data and reference data), Data quality score (Healthy 100), Governance domain (Corporate Functions), Description (data compiled from external sources about company specifics), Use Cases (Monitoring Risk of partners and customers, Reporting with common characteristics and identifiers, Operational data validation), Quality dimensions (Completeness, Quality score Healthy 100), and Glossary terms (Regional Sales, Worldwide Sales, Ecosystem). The "Sales" node in the lineage diagram is also highlighted with a red border.

Roadmap

MS Fabric and Power BI release plan doc

Microsoft Fabric and Power BI release plan documentation

The Microsoft Fabric and Power BI release plan documentation announces the latest updates and timelines to customers as features are prepared for future releases.

Get started

- [OVERVIEW](#)
- What's planned in Microsoft Fabric?
- Admin and governance
- OneLake
- Fabric shared experiences

Analytics

- [OVERVIEW](#)
- Data Warehouse
- Data Engineering
- Data Science

Real-Time Intelligence

- [OVERVIEW](#)
- Real-Time Intelligence

Data Factory

- [OVERVIEW](#)
- Data Factory

Power BI

- [OVERVIEW](#)
- Power BI

- <https://learn.microsoft.com/en-us/fabric/release-plan/>

Links

Links 1/2

- [FabCon 2025: Fueling tomorrow's AI with new agentic capabilities and security innovations in Fabric | Microsoft Fabric Blog](#)
- [Fabric March 2025 Feature Summary | Microsoft Fabric Blog | Microsoft Fabric](#)
- [Copilot and AI Capabilities will be accessible to all paid SKUs in Microsoft Fabric | Microsoft Fabric Blog | Microsoft Fabric](#)
- [The next evolution of OneLake security \(Preview\) | Microsoft Fabric Blog | Microsoft Fabric](#)
- [Migration Assistant for Fabric Data Warehouse \(Preview\) | Microsoft Fabric Blog | Microsoft Fabric](#)
- [Unlock the power of Real-Time Intelligence in the Era of AI: why Fabric Real-Time Intelligence is a game-changer | Microsoft Fabric Blog | Microsoft Fabric](#)
- [Fabric Data Factory: What's New and Latest Roadmap | Microsoft Fabric Blog | Microsoft Fabric](#)
- [Introducing Autoscale Billing for Spark in Microsoft Fabric | Microsoft Fabric Blog | Microsoft Fabric](#)

Links 2/2

- [Empowering agentic AI by integrating Fabric with Azure AI Foundry | Microsoft Fabric Blog | Microsoft Fabric](#)
- [Enhancing AI productivity in Fabric notebooks with Copilot updates | Microsoft Fabric Blog | Microsoft Fabric](#)
- [Mirroring in Fabric – What's new | Microsoft Fabric Blog | Microsoft Fabric](#)
- [New innovations in Microsoft Purview for protected, AI-ready data | Microsoft Security Blog](#)
- [AI Ready Apps: build RAG Data pipeline from Azure Blob Storage to SQL Database in Microsoft Fabric within minutes | Microsoft Fabric Blog | Microsoft Fabric](#)

Thanks!

Andrea Benedetti

Sr. Cloud Architect, Microsoft

<https://www.linkedin.com/in/abenedetti/>

